

# neural-networks-visualization

November 5, 2020

## 1 Neural Networks Visualization

```
In [1]: import os
        import random

        import torch
        from torch.nn import functional as F
        import torchvision
        import torchvision.transforms as T
        import numpy as np
        from scipy.ndimage.filters import gaussian_filter1d
        import matplotlib.pyplot as plt
        from PIL import Image

        torchvision.models.vgg.model_urls[
            "squeezenet1_1"] = "http://webia.lip6.fr/~douillard/rdfia/squeezeNet1_1-f364aa15.pt"
        os.environ["TORCH_HOME"] = "/tmp/torch"

        %matplotlib inline
        plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
        plt.rcParams['image.interpolation'] = 'nearest'
        plt.rcParams['image.cmap'] = 'viridis'
```

### 1.1 Fonctions et variables utiles

```
In [2]: SQUEEZENET_MEAN = np.array([0.485, 0.456, 0.406], dtype=np.float32)
        SQUEEZENET_STD = np.array([0.229, 0.224, 0.225], dtype=np.float32)

        def preprocess(img, size=224):
            transform = T.Compose([
                T.Resize((size, size)),
                T.ToTensor(),
                T.Normalize(mean=SQUEEZENET_MEAN.tolist(),
                           std=SQUEEZENET_STD.tolist()),
                T.Lambda(lambda x: x[None]),
            ])
            return transform(img)
```

```

def deprocess(img, should_rescale=True):
    transform = T.Compose([
        T.Lambda(lambda x: x[0]),
        T.Normalize(mean=[0, 0, 0], std=(1.0 / SQUEEZENET_STD).tolist()),
        T.Normalize(mean=(-SQUEEZENET_MEAN).tolist(), std=[1, 1, 1]),
        T.Lambda(rescale) if should_rescale else T.Lambda(lambda x: x),
        T.ToPILImage(),
    ])
    return transform(img)

def rescale(x):
    low, high = x.min(), x.max()
    x_rescaled = (x - low) / (high - low)
    return x_rescaled

def blur_image(X_np, sigma=1):
    X_np = gaussian_filter1d(X_np, sigma, axis=2)
    X_np = gaussian_filter1d(X_np, sigma, axis=3)
    return torch.tensor(X_np)

def jitter(X, ox, oy):
    """
    Helper function to randomly jitter an image.

    Inputs
    - X: PyTorch Tensor of shape (N, C, H, W)
    - ox, oy: Integers giving number of pixels to jitter along W and H axes

    Returns: A new PyTorch Tensor of shape (N, C, H, W)
    """
    if ox != 0:
        left = X[:, :, :, :-ox]
        right = X[:, :, :, -ox:]
        X = torch.cat([right, left], dim=3)
    if oy != 0:
        top = X[:, :, :-oy]
        bottom = X[:, :, -oy:]
        X = torch.cat([bottom, top], dim=2)
    return X

```

## 1.2 Chargement du modèle

On utilisera le modèle SqueezeNet qui est une modèle léger pré-appris sur ImageNet. Ce modèle sera figé puisque le but n'est pas de modifier/apprendre ses poids mais de les étudier.

In [3]: # Chargement du modèle  
model = torchvision.models.squeezenet1\_1(pretrained=True)

```

# Modele en mode test
model.eval()

# Freeze les poids
for param in model.parameters():
    param.requires_grad = False

```

### 1.3 Chargement d'images d'exemples

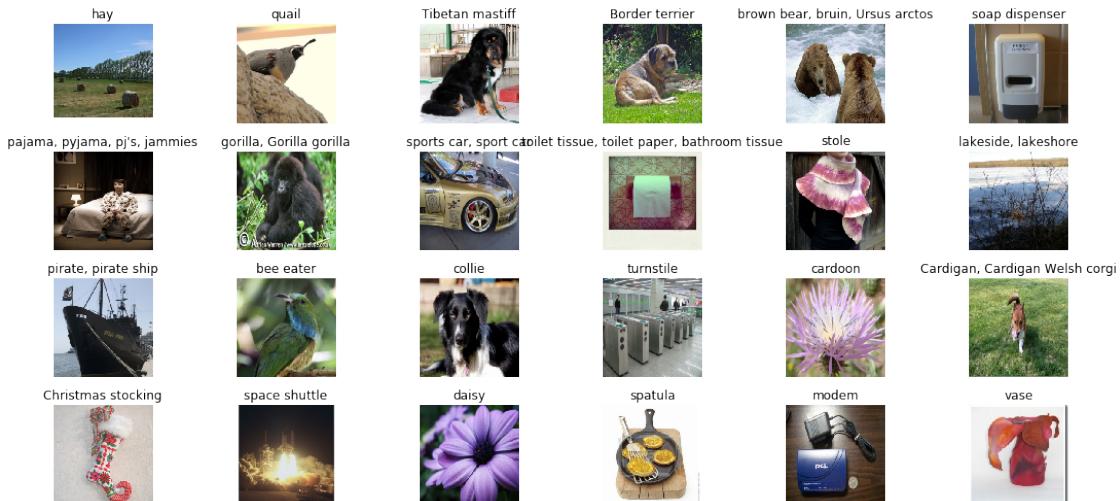
Permet de chargement dans les variables `X`, `y`, `class_names` 25 exemples de l'ensemble de validation d'ImageNet. `X` contient les images, `y` l'indice de la classe de chaque images, et `class_names` un dictionnaire permet d'obtenir le nom d'une classe à partir de son numéro.

```
In [4]: f = np.load("imagenet_val_25.npz", allow_pickle=True)
X, y, class_names = f["X"], f["y"], f["label_map"].item()
```

```

plt.figure(figsize=(15, 7))
for i in range(24):
    plt.subplot(4, 6, i + 1)
    plt.imshow(X[i])
    plt.title(class_names[y[i]])
    plt.axis('off')
plt.gcf().tight_layout()

```



```
In [5]: type(X), type(y), type(class_names)
```

```
Out[5]: (numpy.ndarray, numpy.ndarray, dict)
```

```
In [6]: len(X), len(y), len(class_names)
```

```
Out[6]: (25, 25, 999)
```

## 2 Saliency Maps

**Conseil :** Pour sélectionner 1 valeur particulière pour chaque ligne d'une matrice, vous pouvez faire comme cela :

```
x = torch.Tensor([[0.1, 0.0, 0.5, 0.1, 0.1],  
                  [0.0, 0.1, 0.0, 0.6, 0.2],  
                  [0.7, 0.1, 0.1, 0.3, 0.0]])  
x[np.arange(3), [2, 3, 0]]  
# 0.5000  
# 0.6000  
# 0.7000  
#[torch.FloatTensor of size 3]  
  
In [7]: def compute_saliency_maps(X, y, model):  
    """  
    Compute a class saliency map using the model for images X and labels y.  
  
    Input:  
    - X: Input images; Tensor of shape (N, 3, H, W)  
    - y: Labels for X; LongTensor of shape (N, )  
    - model: A pretrained CNN that will be used to compute the saliency map.  
  
    Returns:  
    - saliency: A Tensor of shape (N, H, W) giving the saliency maps for the input  
    images.  
    """  
    X.requires_grad=True  
  
    #####  
    # TODO: Implement this function. Perform a forward and backward pass through #  
    # the model to compute the gradient of the correct class score with respect #  
    # to each input image. You first want to compute the loss over the correct #  
    # scores, and then compute the gradients with a backward pass. #  
    #####  
  
    yo_batch = model.forward(X)  
  
    yo_batch = yo_batch[np.arange(len(y)), y]  
  
    (yo_batch.sum()).backward()  
  
    saliency = torch.abs(X.grad.data)  
    saliency = torch.max(saliency, dim=1)[0]  
  
    #####  
    # END OF YOUR CODE  
    #####  
    return saliency
```

Testez votre code avec la fonction ci-dessous :

In [8]: model

Out[8]: SqueezeNet(  
    (features): Sequential(  
        (0): Conv2d(3, 64, kernel\_size=(3, 3), stride=(2, 2))  
        (1): ReLU(inplace=True)  
        (2): MaxPool2d(kernel\_size=3, stride=2, padding=0, dilation=1, ceil\_mode=True)  
        (3): Fire(  
            (squeeze): Conv2d(64, 16, kernel\_size=(1, 1), stride=(1, 1))  
            (squeeze\_activation): ReLU(inplace=True)  
            (expand1x1): Conv2d(16, 64, kernel\_size=(1, 1), stride=(1, 1))  
            (expand1x1\_activation): ReLU(inplace=True)  
            (expand3x3): Conv2d(16, 64, kernel\_size=(3, 3), stride=(1, 1), padding=(1, 1))  
            (expand3x3\_activation): ReLU(inplace=True)  
        )  
        (4): Fire(  
            (squeeze): Conv2d(128, 16, kernel\_size=(1, 1), stride=(1, 1))  
            (squeeze\_activation): ReLU(inplace=True)  
            (expand1x1): Conv2d(16, 64, kernel\_size=(1, 1), stride=(1, 1))  
            (expand1x1\_activation): ReLU(inplace=True)  
            (expand3x3): Conv2d(16, 64, kernel\_size=(3, 3), stride=(1, 1), padding=(1, 1))  
            (expand3x3\_activation): ReLU(inplace=True)  
        )  
        (5): MaxPool2d(kernel\_size=3, stride=2, padding=0, dilation=1, ceil\_mode=True)  
        (6): Fire(  
            (squeeze): Conv2d(128, 32, kernel\_size=(1, 1), stride=(1, 1))  
            (squeeze\_activation): ReLU(inplace=True)  
            (expand1x1): Conv2d(32, 128, kernel\_size=(1, 1), stride=(1, 1))  
            (expand1x1\_activation): ReLU(inplace=True)  
            (expand3x3): Conv2d(32, 128, kernel\_size=(3, 3), stride=(1, 1), padding=(1, 1))  
            (expand3x3\_activation): ReLU(inplace=True)  
        )  
        (7): Fire(  
            (squeeze): Conv2d(256, 32, kernel\_size=(1, 1), stride=(1, 1))  
            (squeeze\_activation): ReLU(inplace=True)  
            (expand1x1): Conv2d(32, 128, kernel\_size=(1, 1), stride=(1, 1))  
            (expand1x1\_activation): ReLU(inplace=True)  
            (expand3x3): Conv2d(32, 128, kernel\_size=(3, 3), stride=(1, 1), padding=(1, 1))  
            (expand3x3\_activation): ReLU(inplace=True)  
        )  
        (8): MaxPool2d(kernel\_size=3, stride=2, padding=0, dilation=1, ceil\_mode=True)  
        (9): Fire(  
            (squeeze): Conv2d(256, 48, kernel\_size=(1, 1), stride=(1, 1))  
            (squeeze\_activation): ReLU(inplace=True)  
            (expand1x1): Conv2d(48, 192, kernel\_size=(1, 1), stride=(1, 1))  
            (expand1x1\_activation): ReLU(inplace=True)

```

(expand3x3): Conv2d(48, 192, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(expand3x3_activation): ReLU(inplace=True)
)
(10): Fire(
    (squeeze): Conv2d(384, 48, kernel_size=(1, 1), stride=(1, 1))
    (squeeze_activation): ReLU(inplace=True)
    (expand1x1): Conv2d(48, 192, kernel_size=(1, 1), stride=(1, 1))
    (expand1x1_activation): ReLU(inplace=True)
    (expand3x3): Conv2d(48, 192, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (expand3x3_activation): ReLU(inplace=True)
)
(11): Fire(
    (squeeze): Conv2d(384, 64, kernel_size=(1, 1), stride=(1, 1))
    (squeeze_activation): ReLU(inplace=True)
    (expand1x1): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1))
    (expand1x1_activation): ReLU(inplace=True)
    (expand3x3): Conv2d(64, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (expand3x3_activation): ReLU(inplace=True)
)
(12): Fire(
    (squeeze): Conv2d(512, 64, kernel_size=(1, 1), stride=(1, 1))
    (squeeze_activation): ReLU(inplace=True)
    (expand1x1): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1))
    (expand1x1_activation): ReLU(inplace=True)
    (expand3x3): Conv2d(64, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (expand3x3_activation): ReLU(inplace=True)
)
)
(classifier): Sequential(
    (0): Dropout(p=0.5, inplace=False)
    (1): Conv2d(512, 1000, kernel_size=(1, 1), stride=(1, 1))
    (2): ReLU(inplace=True)
    (3): AdaptiveAvgPool2d(output_size=(1, 1))
)
)

```

```

In [9]: def show_saliency_maps(X, y, model):
    # Convert X and y from numpy arrays to Torch Tensors
    X_tensor = torch.cat([preprocess(Image.fromarray(x)) for x in X], dim=0)
    y_tensor = torch.LongTensor(y)

    # Compute saliency maps for images in X
    saliency = compute_saliency_maps(X_tensor, y_tensor, model)

    # Convert the saliency map from Torch Tensor to numpy array and show images
    # and saliency maps together.
    saliency = saliency.numpy()
    N = X.shape[0]

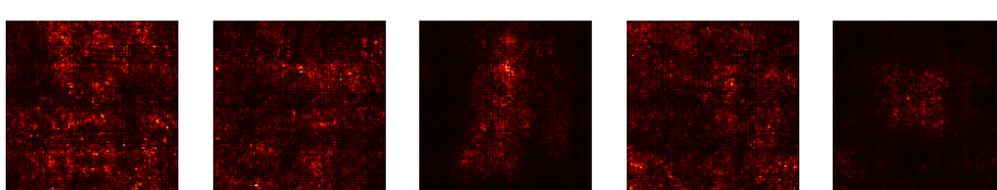
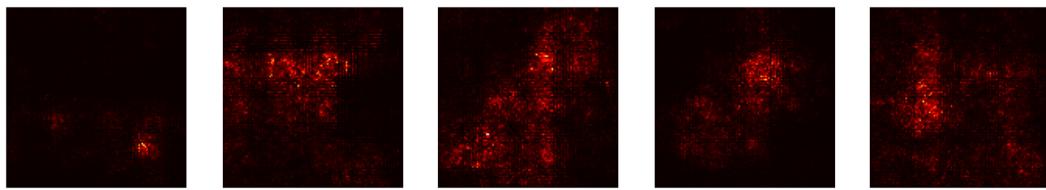
```

```

for i in range(N):
    plt.subplot(2, N, i + 1)
    plt.imshow(X[i])
    plt.axis('off')
    plt.title(class_names[y[i]])
    plt.subplot(2, N, N + i + 1)
    #print(type(np.transpose(saliency[i], (1,2,0))))
    plt.imshow(saliency[i], cmap=plt.cm.hot)
    plt.axis('off')
    plt.gcf().set_size_inches(12, 5)
plt.show()

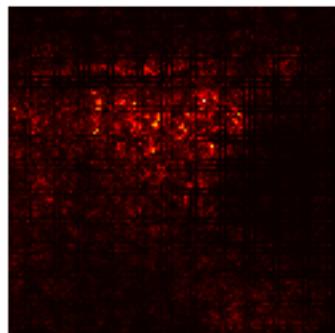
for i in range(2): # range(5) pour tester toutes les images
    show_saliency_maps(X[5*i:5*i+5], y[5*i:5*i+5], model)

```



In [10]: `show_saliency_maps(X[1:1+1], y[1:1+1], model)`

quail



### 3 Fooling Images

Code pour calculer une image de facon à ce qu'elle soit classée dans une classe target\_y autre que la classe réelle (en modifiant l'image et pas les poids du réseau).

Les 2 premiers blocs permettent de faire des tests de facon interractive pour écrire votre code. Une fois que votre code semble fonctionner, compléter la fonction dans le 3e bloc et tester sur differentes images avec le 4e bloc.

```
In [11]: # Init du test
X_tensor = torch.Tensor(preprocess(Image.fromarray(X[0])))
target_y = 6
X_fooling = X_tensor.clone()
X_fooling.requires_grad = True
learning_rate = 1.0
```

```
In [12]: t = 0
while True:
    yo = model.forward(X_fooling)
    yy = yo.argmax(dim=1)[0].item()
    if yy == target_y:
        break
```

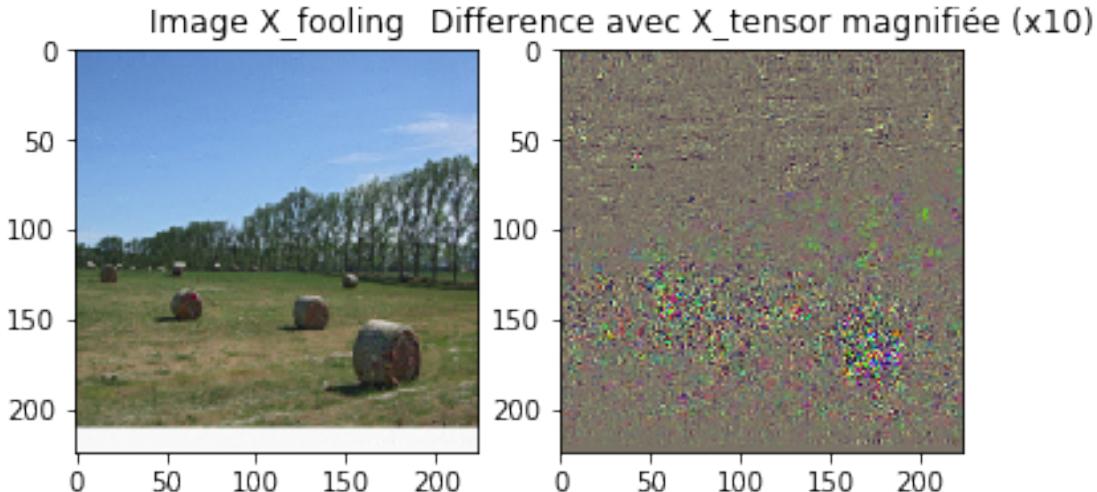
```

yo[0,target_y].backward()
with torch.no_grad():
    X_fooling.data += learning_rate*X_fooling.grad.data
print(t,end='->')
t+=1

# Affichage de l'image X_fooling et ses modifs
plt.subplot(1, 2, 1)
plt.imshow(np.asarray(deprocess(X_fooling.clone())).astype(np.uint8))
plt.title("Image X_fooling")
plt.subplot(1, 2, 2)
plt.imshow(np.asarray(deprocess(10* (X_fooling - X_tensor), should_rescale=False)))
plt.title("Difference avec X_tensor magnifiée (x10)")
plt.show()

```

0->1->2->



```

In [13]: def make_fooling_image(X, target_y, model, nb_iterations=100):
    """
    Generate a fooling image that is close to X, but that the model classifies
    as target_y.

    Inputs:
    - X: Input image; Tensor of shape (1, 3, 224, 224)
    - target_y: An integer in the range [0, 1000)
    - model: A pretrained CNN

    Returns:
    - X_fooling: An image that is close to X, but that is classified as target_y
    by the model.
    """

```

```

"""
# Initialize our fooling image to the input image, and wrap it in a Variable.
X_fooling = X.clone()
X_fooling.requires_grad = True

learning_rate = 0.1
#####
# TODO: Generate a fooling image X_fooling that the model will classify as #
# the class target_y. You should perform gradient ascent on the score of the #
# target class, stopping when the model is fooled.                                #
# When computing an update step, first normalize the gradient:                   #
#   dX = learning_rate * g / ||g||_2                                           #
# #
# You should write a training loop.                                              #
# #
# HINT: For most examples, you should be able to generate a fooling image     #
# in fewer than 100 iterations of gradient ascent.                               #
# You can print your progress over iterations to check your algorithm.        #
#####

while True:
    yo = model.forward(X_fooling)
    yy = yo.argmax(dim=1)[0].item()
    if yy == target_y:
        break
    yo[0,target_y].backward()
    with torch.no_grad():
        X_fooling.data += learning_rate*X_fooling.grad.data

#####
# END OF YOUR CODE
#####
return X_fooling

```

```

In [14]: # Indice de l'image à modifier et de la classe cible
idx = 1
target_y = 6

# Préparation du tenseur X et sa version "fooling"
X_tensor = torch.cat([preprocess(Image.fromarray(x)) for x in X], dim=0)
X_fooling = make_fooling_image(X_tensor[idx:idx+1], target_y, model)

# vérification de la classe prédite
scores = model(X_fooling)
assert target_y == scores.data.max(1)[1][0], 'The model is not fooled!'

# Affichage
X_fooling_np = deprocess(X_fooling.clone())

```

```

X_fooling_np = np.asarray(X_fooling_np).astype(np.uint8)

plt.subplot(1, 4, 1)
plt.imshow(X[idx])
plt.title(class_names[y[idx]])
plt.axis('off')

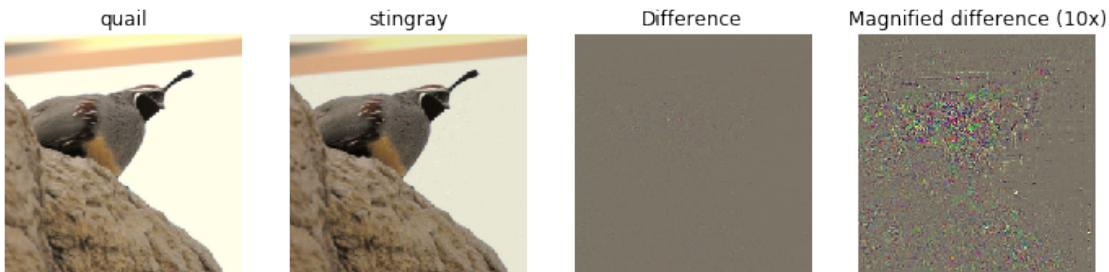
plt.subplot(1, 4, 2)
plt.imshow(X_fooling_np)
plt.title(class_names[target_y])
plt.axis('off')

plt.subplot(1, 4, 3)
X_pre = preprocess(Image.fromarray(X[idx]))
diff = np.asarray(deprocess(X_fooling - X_pre, should_rescale=False))
plt.imshow(diff)
plt.title('Difference')
plt.axis('off')

plt.subplot(1, 4, 4)
diff = np.asarray(deprocess(10 * (X_fooling - X_pre), should_rescale=False))
plt.imshow(diff)
plt.title('Magnified difference (10x)')
plt.axis('off')

plt.gcf().set_size_inches(12, 5)
plt.show()

```



In [15]: X\_fooling.shape, np.array([target\_y])[None].shape

Out[15]: (torch.Size([1, 3, 224, 224]), (1, 1))

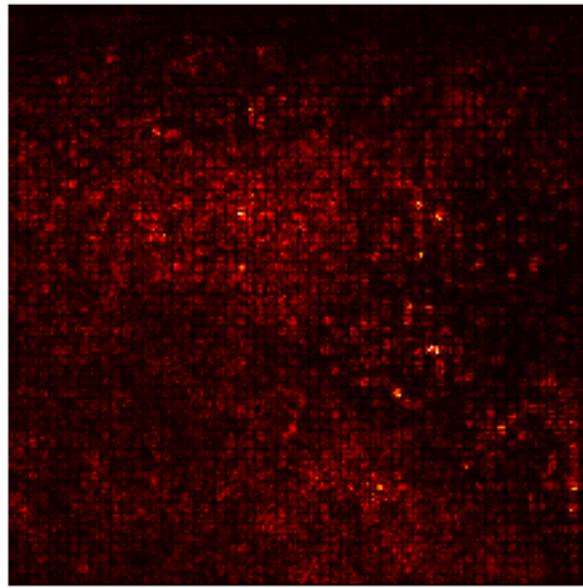
In [16]: type(compute\_saliency\_maps(X\_fooling, np.array([target\_y]), model))

Out[16]: torch.Tensor

In [17]: plt.imshow(compute\_saliency\_maps(X\_fooling, np.array([target\_y]), model)[0] , cmap=plt.cm.viridis)
plt.title("saliency map for fooling image \n stingray class \n quail --> stingray")
plt.axis("off")

```
Out[17]: (-0.5, 223.5, 223.5, -0.5)
```

saliency map for fooling image  
stingray class  
quail --> stingray



## 4 Class visualization

Code permettant de calculer une image maximisant le score d'une classe, sujet à un certain nombre de régularisations.

```
In [18]: def create_class_visualization(target_y, model, dtype, init_img=None, l2_reg=1e-3, learn_rate=0.01, num_iterations=200, blur_every=10, max_jitter=16, show=False):
```

*Generate an image to maximize the score of target\_y under a pretrained model.*

*Inputs:*

- *target\_y: Integer in the range [0, 1000) giving the index of the class*
- *model: A pretrained CNN that will be used to generate the image*
- *dtype: Torch datatype to use for computations*

*Keyword arguments:*

- *init\_img: Initial image to use (if None, will be random)*
- *l2\_reg: Strength of L2 regularization on the image*
- *learning\_rate: How big of a step to take*
- *num\_iterations: How many iterations to use*
- *blur\_every: How often to blur the image as an implicit regularizer*

```

- max_jitter: How much to gjitter the image as an implicit regularizer
- show_every: How often to show the intermediate result
"""

model.type(dtype)

# Randomly initialize the image as a PyTorch Tensor:
if init_img is None:
    img = torch.randn(1, 3, 224, 224).mul_(1.0).type(dtype)
else:
    img = init_img.clone().mul_(1.0).type(dtype)
img.requires_grad = True

for t in range(num_iterations):
    # Randomly jitter the image a bit; this gives slightly nicer results
    ox, oy = random.randint(0, max_jitter), random.randint(0, max_jitter)
    img.data.copy_(jitter(img.data, ox, oy))

    ##### TODO: Use the model to compute the gradient of the score for the
    # class target_y with respect to the pixels of the image, and make a
    # gradient step on the image using the learning rate. Don't forget the
    # L2 regularization term!
    # Be very careful about the signs of elements in your code.
    #####
    y_hat = model.forward(img)
    #print(y_hat.argmax(dim=1)[0].item(), class_names[y_hat.argmax(dim=1)[0].item()])
    loss = y_hat[0, target_y] - 12_reg * torch.norm(img, 2)
    loss.backward()

    with torch.no_grad():
        img += learning_rate*img.grad

    ##### END OF YOUR CODE #####
    img.data.copy_(jitter(img.data, -ox, -oy))

    # As regularizer, clamp and periodically blur the image
    for c in range(3):
        lo = float(-SQUEEZENET_MEAN[c] / SQUEEZENET_STD[c])
        hi = float((1.0 - SQUEEZENET_MEAN[c]) / SQUEEZENET_STD[c])
        img.data[:, c].clamp_(min=lo, max=hi)
    if t % blur_every == 0:
        img.data = blur_image(img.data.cpu().numpy(), sigma=0.5)

```

```

# Periodically show the image
if t == 0 or (t + 1) % show_every == 0 or t == num_iterations - 1:
    plt.imshow(deprocess(img.clone().cpu()))
    class_name = class_names[target_y]
    cn = class_names[y_hat.argmax(dim=1)[0].item()]
    plt.title('y_hat:%s\ny:%s\nIteration %d / %d' % (cn, class_name, t + 1, num_iterations))
    plt.gcf().set_size_inches(4, 4)
    plt.axis('off')
    plt.show()

return deprocess(img.cpu())

```

Test avec diverses classes en partant d'un bruit aléatoire :

```

In [19]: dtype = torch.FloatTensor
        # dtype = torch.cuda.FloatTensor # Uncomment this to use GPU
        model.type(dtype)

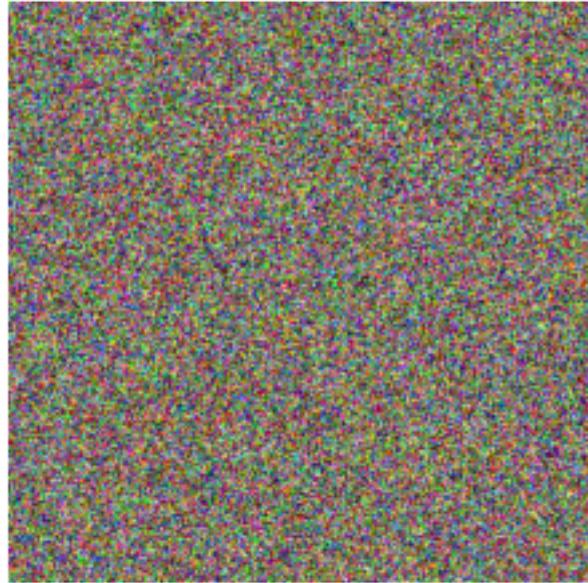
l2_reg = 1e-3
learning_rate = 0.1

num_iterations = 1000

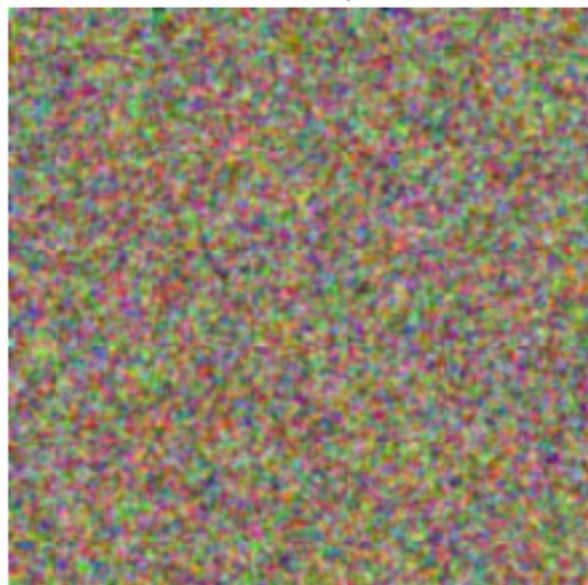
#target_y = 281 # Tabby cat
# target_y = 187 # Yorkshire Terrier
#target_y = 76 # Tarantula
# target_y = 78 # Tick
# target_y = 683 # Oboe
target_y = 366 # Gorilla
# target_y = 604 # Hourglass
# target_y = np.random.randint(1000) # Classe aléatoire
out = create_class_visualization(target_y, model, dtype, show_every=50, num_iterations=num_iterations)

```

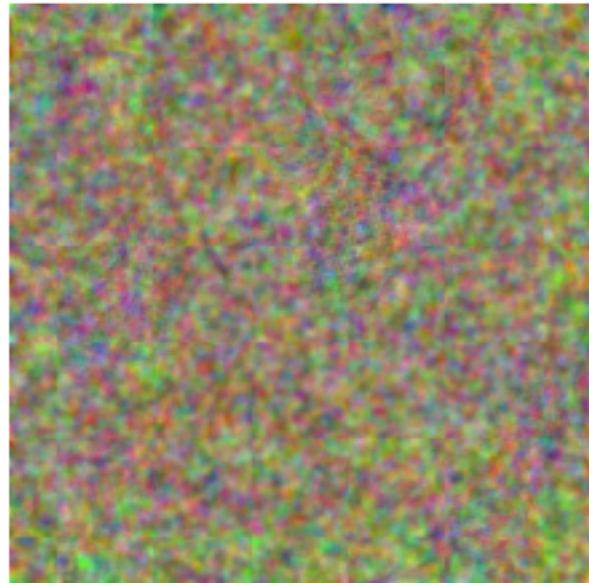
y\_hat:doormat, welcome mat  
y:gorilla, Gorilla gorilla  
Iteration 1 / 1000



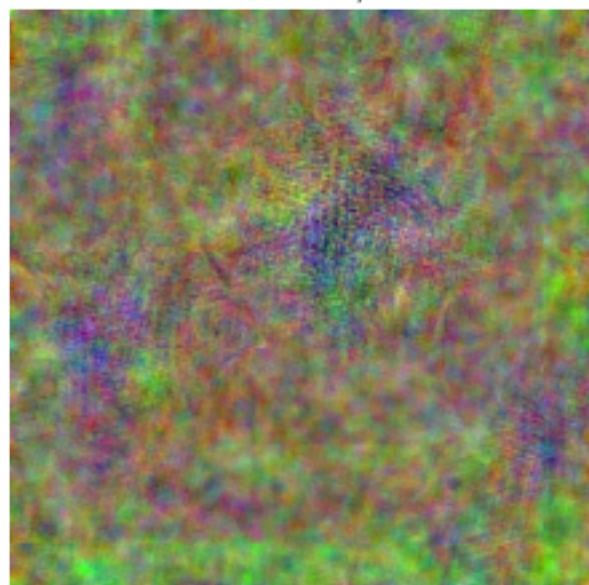
y\_hat:skunk, polecat, wood pussy  
y:gorilla, Gorilla gorilla  
Iteration 50 / 1000



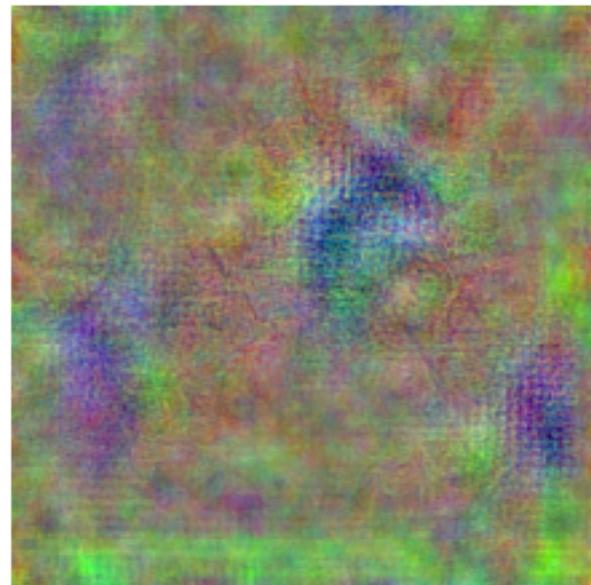
y\_hat:bustard  
y:gorilla, Gorilla gorilla  
Iteration 100 / 1000



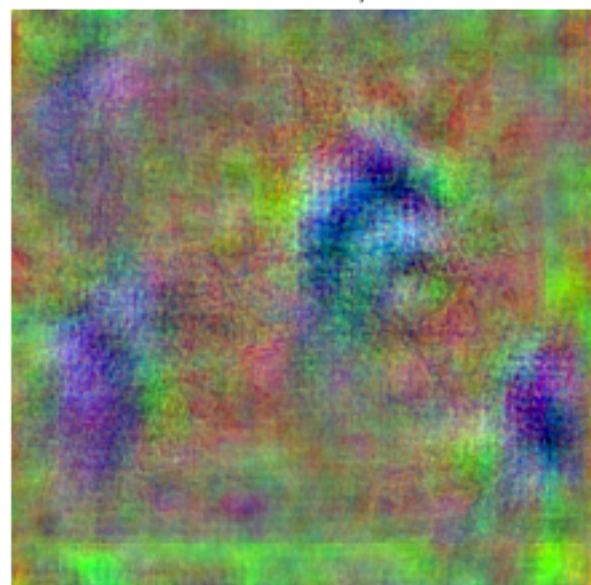
y\_hat:grey fox, gray fox, Urocyon cinereoargenteus  
y:gorilla, Gorilla gorilla  
Iteration 150 / 1000



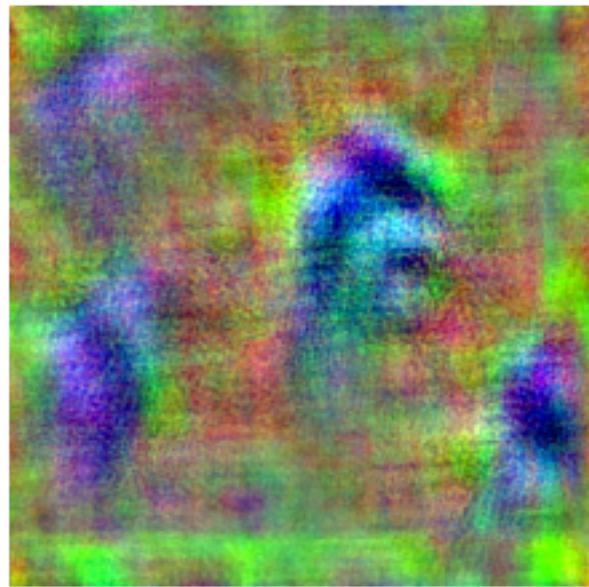
y\_hat:peacock  
y:gorilla, Gorilla gorilla  
Iteration 200 / 1000



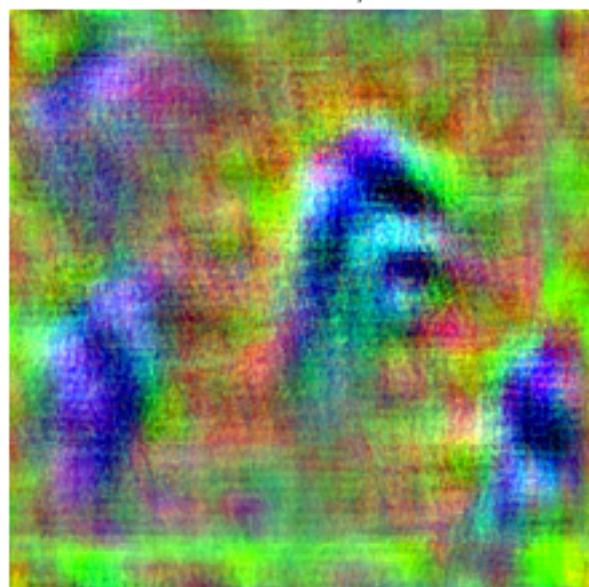
y\_hat:peacock  
y:gorilla, Gorilla gorilla  
Iteration 250 / 1000



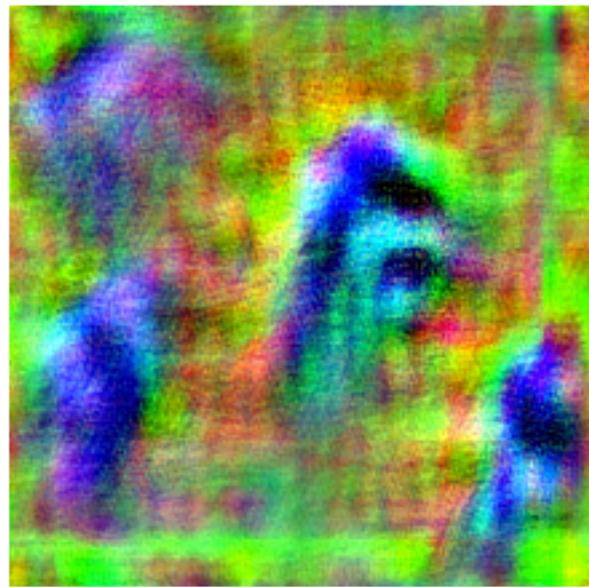
y\_hat:gorilla, Gorilla gorilla  
y:gorilla, Gorilla gorilla  
Iteration 300 / 1000



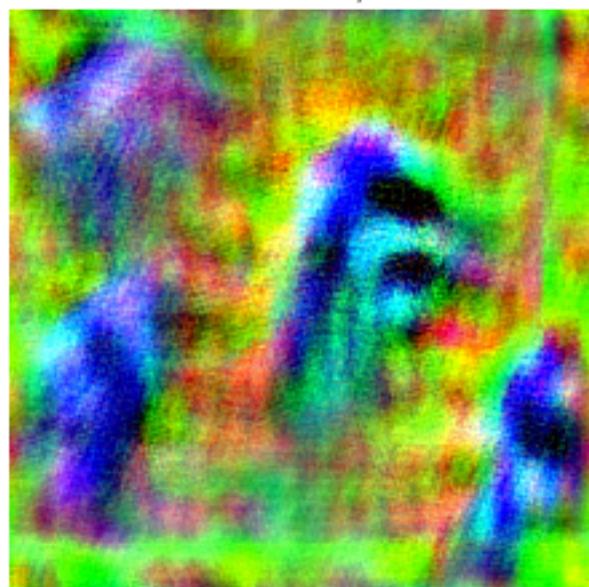
y\_hat:gorilla, Gorilla gorilla  
y:gorilla, Gorilla gorilla  
Iteration 350 / 1000



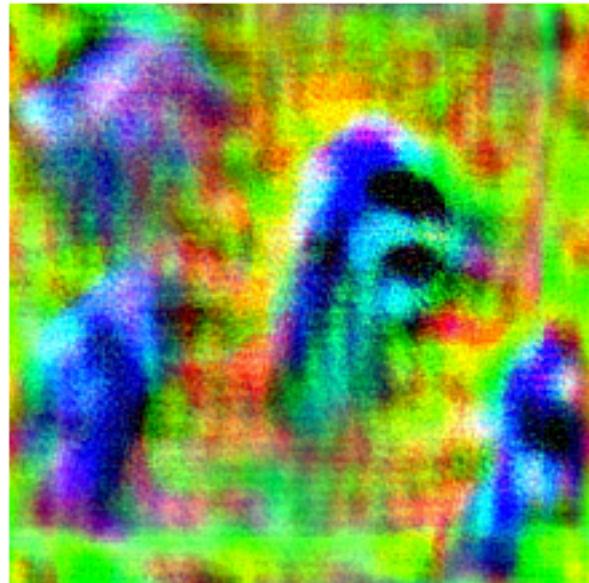
y\_hat:gorilla, Gorilla gorilla  
y:gorilla, Gorilla gorilla  
Iteration 400 / 1000



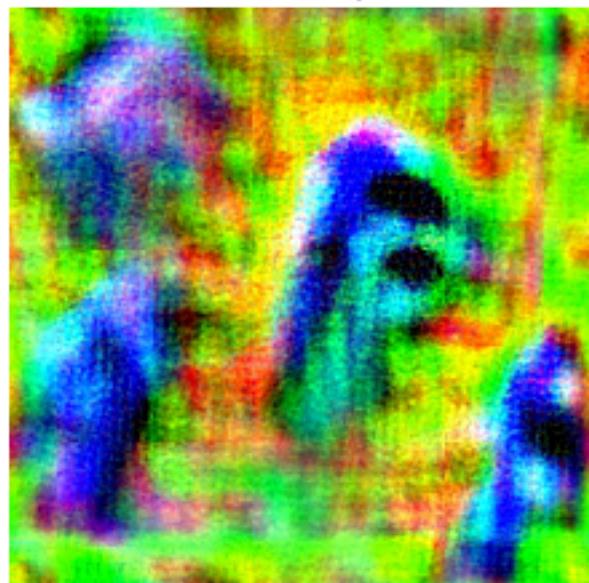
y\_hat:gorilla, Gorilla gorilla  
y:gorilla, Gorilla gorilla  
Iteration 450 / 1000



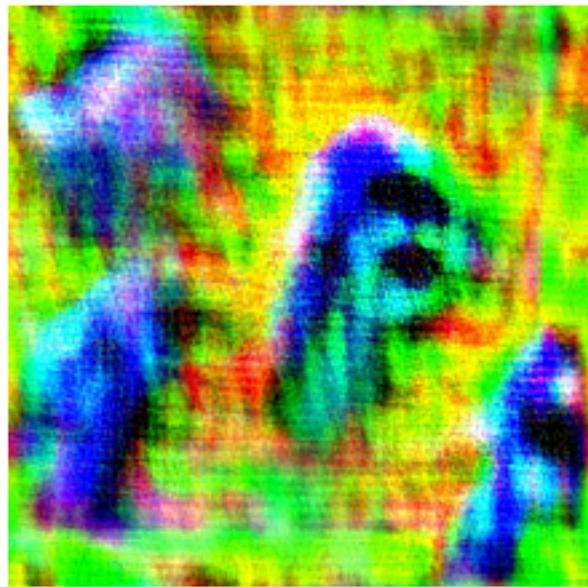
y\_hat:gorilla, Gorilla gorilla  
y:gorilla, Gorilla gorilla  
Iteration 500 / 1000



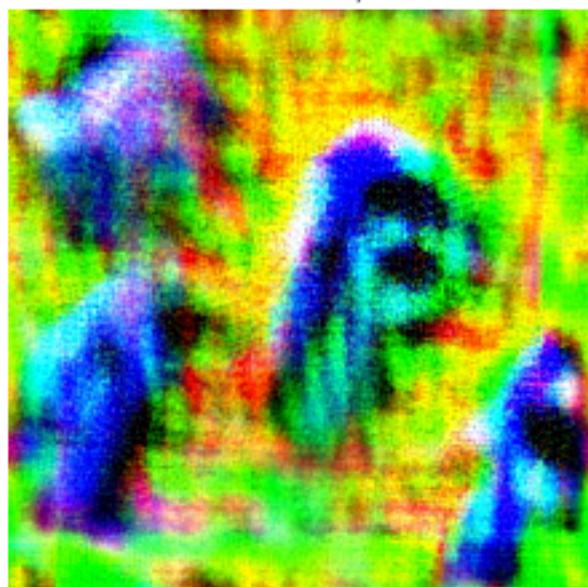
y\_hat:gorilla, Gorilla gorilla  
y:gorilla, Gorilla gorilla  
Iteration 550 / 1000



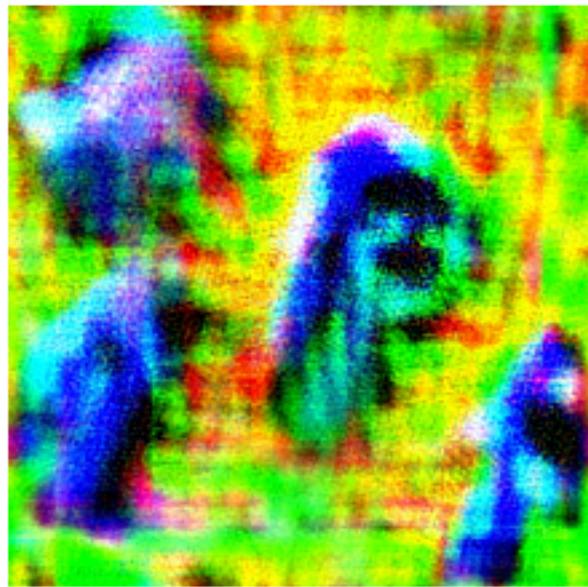
y\_hat:gorilla, Gorilla gorilla  
y:gorilla, Gorilla gorilla  
Iteration 600 / 1000



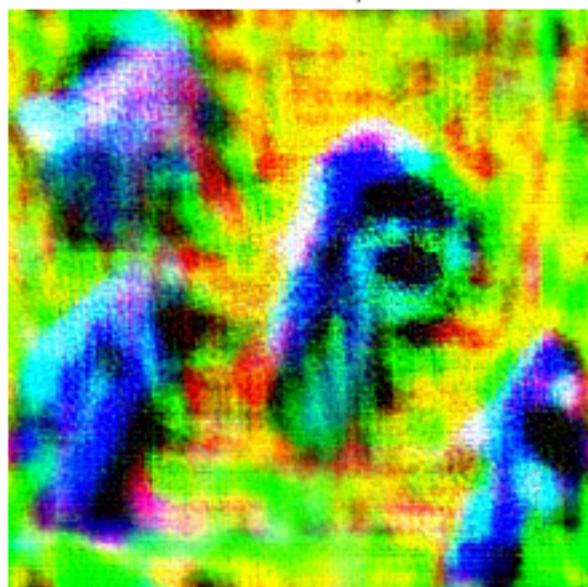
y\_hat:gorilla, Gorilla gorilla  
y:gorilla, Gorilla gorilla  
Iteration 650 / 1000



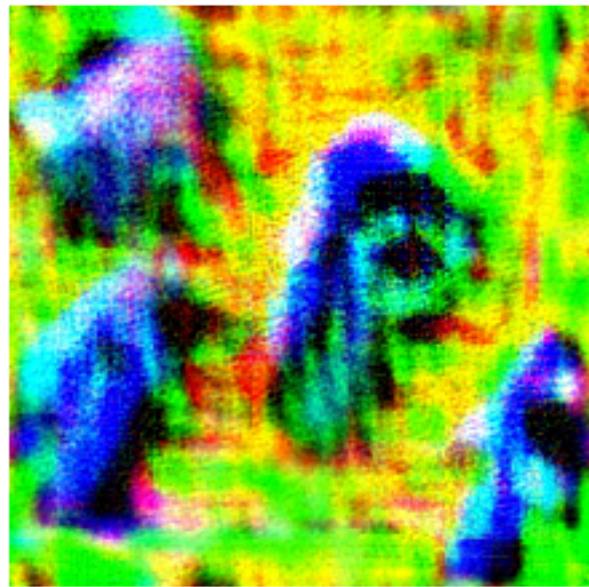
y\_hat:gorilla, Gorilla gorilla  
y:gorilla, Gorilla gorilla  
Iteration 700 / 1000



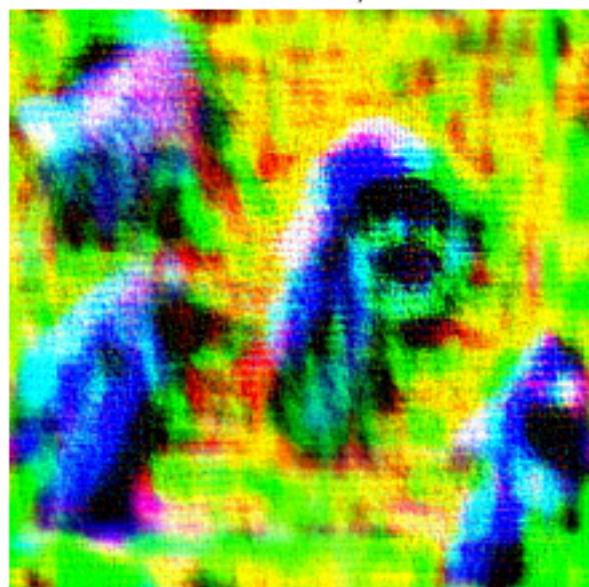
y\_hat:gorilla, Gorilla gorilla  
y:gorilla, Gorilla gorilla  
Iteration 750 / 1000



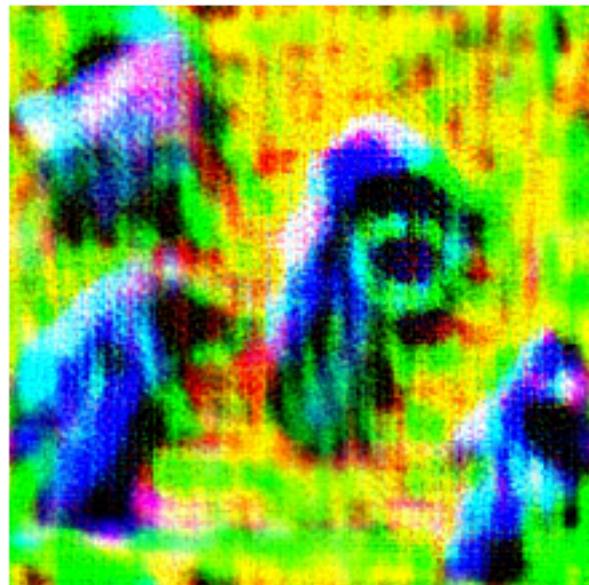
y\_hat:gorilla, Gorilla gorilla  
y:gorilla, Gorilla gorilla  
Iteration 800 / 1000



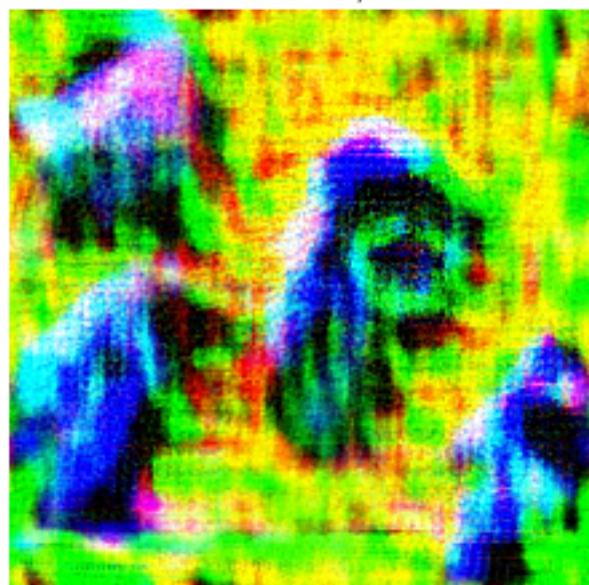
y\_hat:gorilla, Gorilla gorilla  
y:gorilla, Gorilla gorilla  
Iteration 850 / 1000



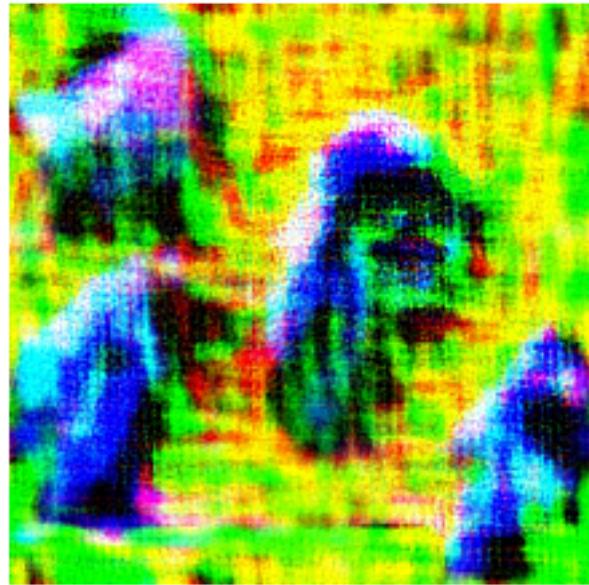
y\_hat:gorilla, Gorilla gorilla  
y:gorilla, Gorilla gorilla  
Iteration 900 / 1000



y\_hat:gorilla, Gorilla gorilla  
y:gorilla, Gorilla gorilla  
Iteration 950 / 1000

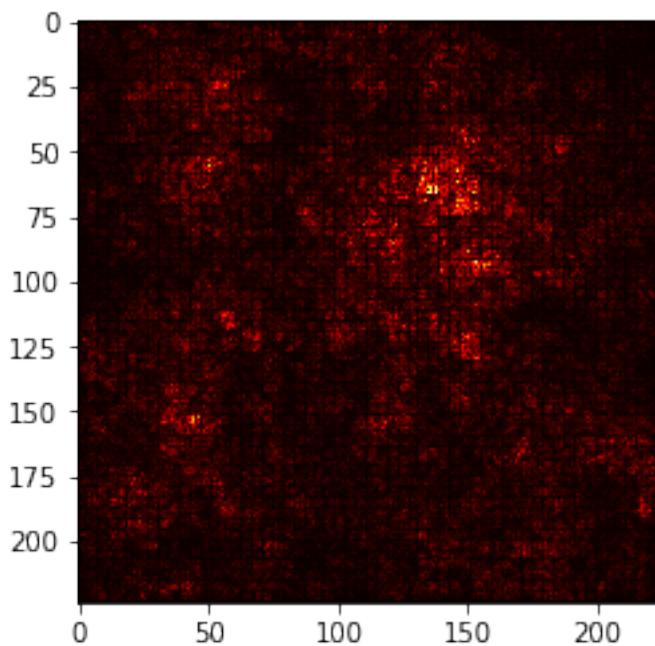


y\_hat:groenendael  
y:gorilla, Gorilla gorilla  
Iteration 1000 / 1000



In [20]: plt.imshow(compute\_saliency\_maps(preprocess(out), np.array([target\_y]), model)[0] , c

Out[20]: <matplotlib.image.AxesImage at 0x7fc8b44b5b00>



Test en partant d'une image d'ImageNet :

```
In [21]: # Init du test
    img_ind = 4

    target_y = y[img_ind]

    #target_y = 281 # Tabby cat

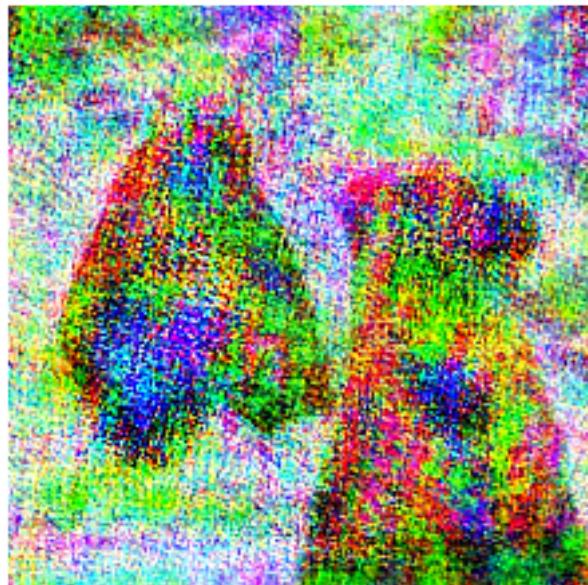
    l2_reg = 1e-2
    learning_rate = 0.1

    X_tensor = torch.Tensor(preprocess(Image.fromarray(X[img_ind])))
    out = create_class_visualization(target_y, model, dtype, init_img=X_tensor, show_every=100)

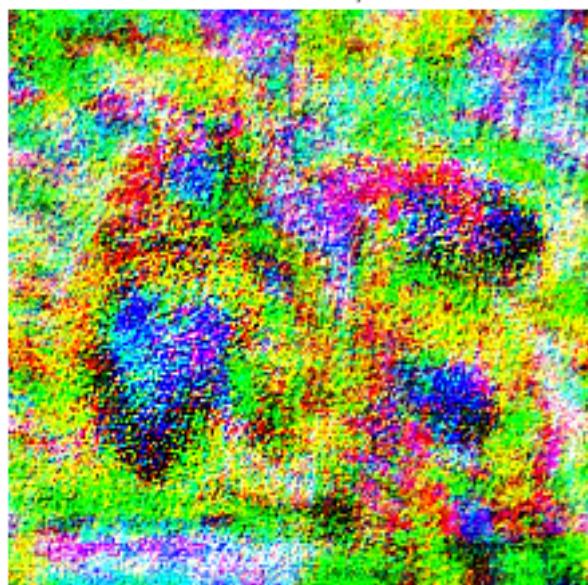
    y_hat: brown bear, bruin, Ursus arctos
    y:brown bear, bruin, Ursus arctos
    Iteration 1 / 1000
```



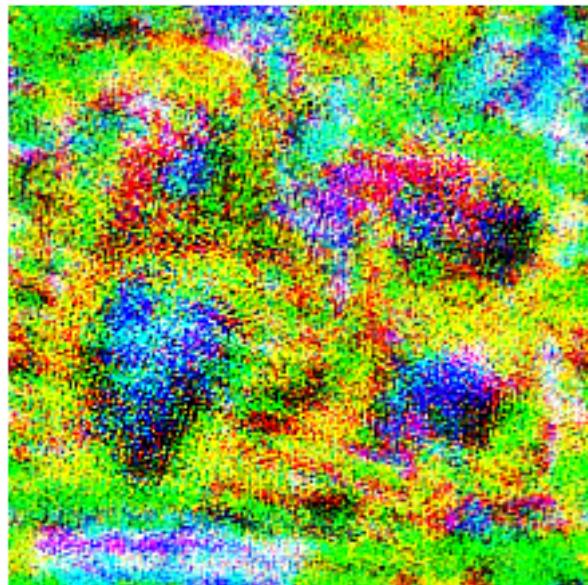
y\_hat:coral reef  
y:brown bear, bruin, Ursus arctos  
Iteration 50 / 1000



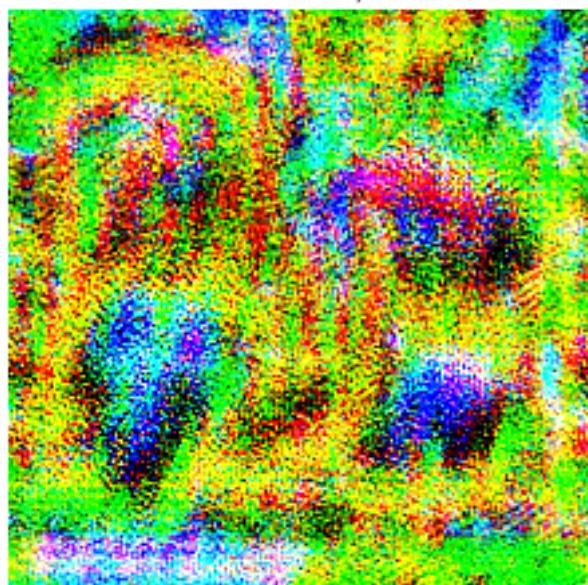
y\_hat:jigsaw puzzle  
y:brown bear, bruin, Ursus arctos  
Iteration 100 / 1000



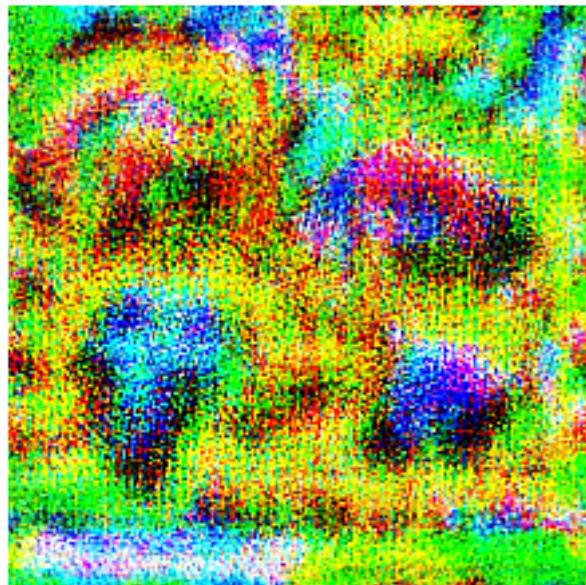
y\_hat:jigsaw puzzle  
y:brown bear, bruin, Ursus arctos  
Iteration 150 / 1000



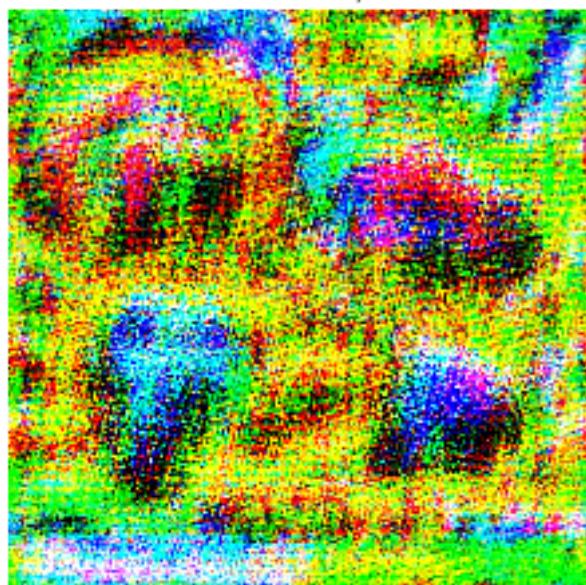
y\_hat:peacock  
y:brown bear, bruin, Ursus arctos  
Iteration 200 / 1000



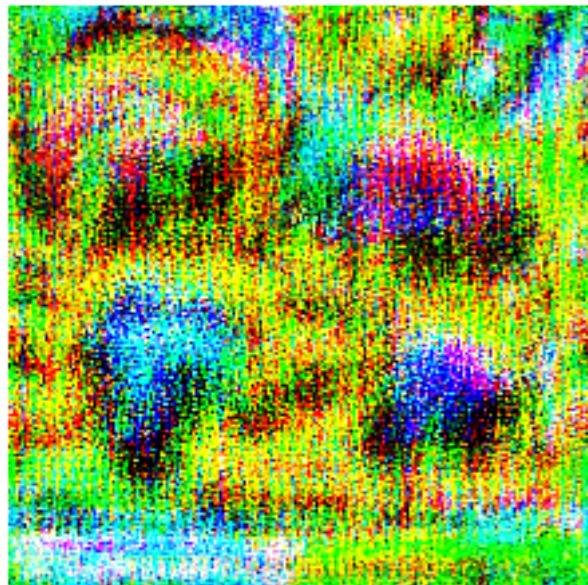
y\_hat:jigsaw puzzle  
y:brown bear, bruin, Ursus arctos  
Iteration 250 / 1000



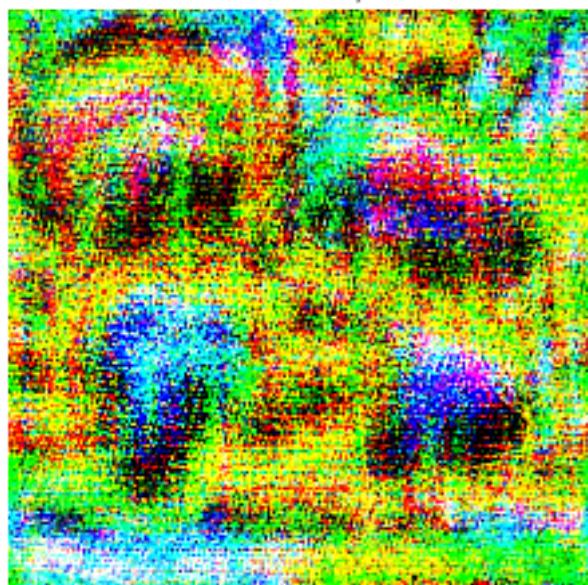
y\_hat:peacock  
y:brown bear, bruin, Ursus arctos  
Iteration 300 / 1000



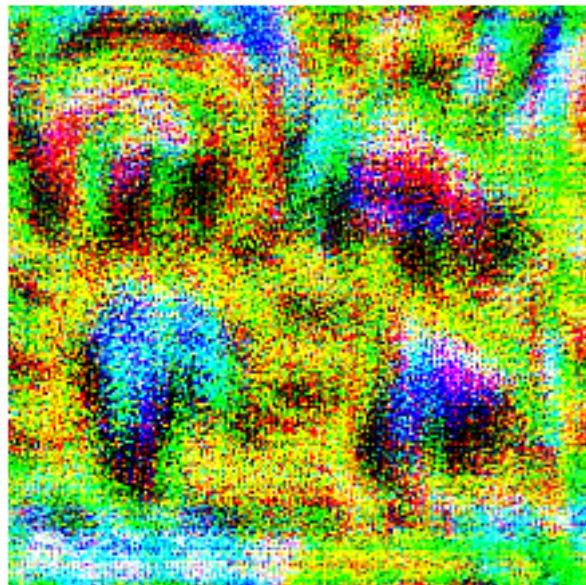
y\_hat:jigsaw puzzle  
y:brown bear, bruin, Ursus arctos  
Iteration 350 / 1000



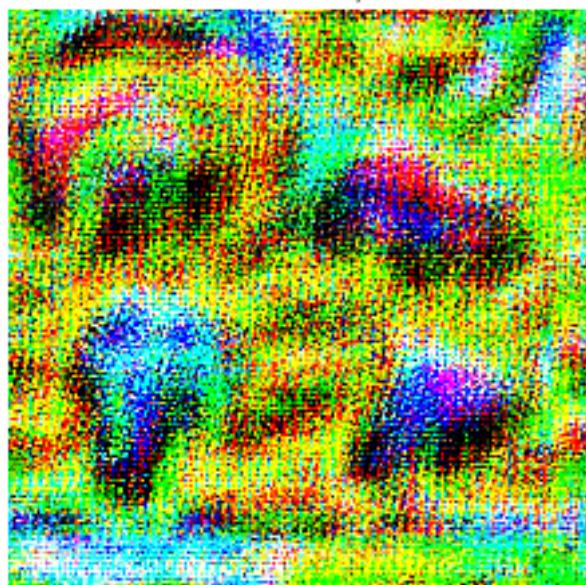
y\_hat:peacock  
y:brown bear, bruin, Ursus arctos  
Iteration 400 / 1000



y\_hat:coral reef  
y:brown bear, bruin, Ursus arctos  
Iteration 450 / 1000



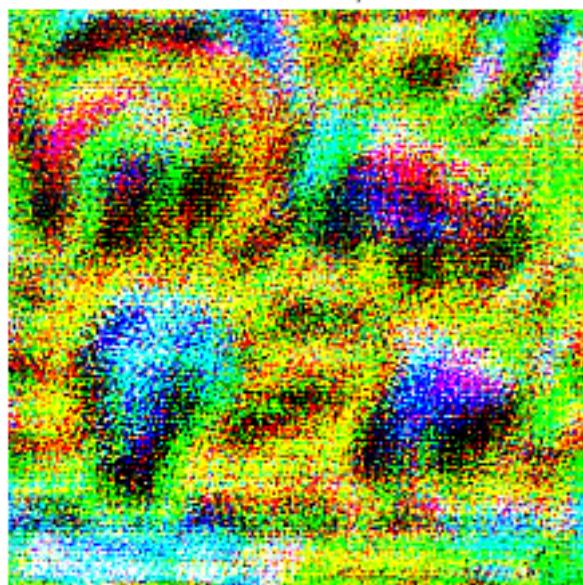
y\_hat:jigsaw puzzle  
y:brown bear, bruin, Ursus arctos  
Iteration 500 / 1000



y\_hat:peacock  
y:brown bear, bruin, Ursus arctos  
Iteration 550 / 1000



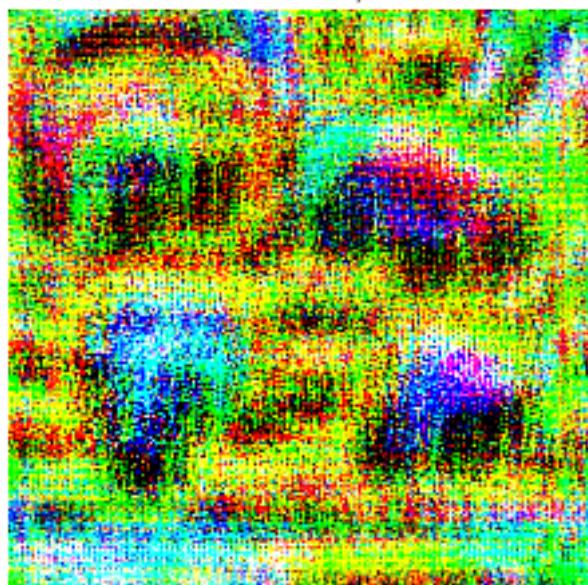
y\_hat:peacock  
y:brown bear, bruin, Ursus arctos  
Iteration 600 / 1000



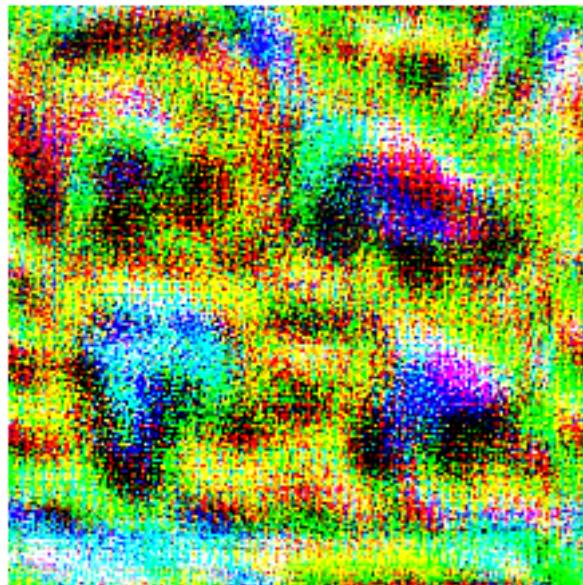
y\_hat:brown bear, bruin, Ursus arctos  
y:brown bear, bruin, Ursus arctos  
Iteration 650 / 1000



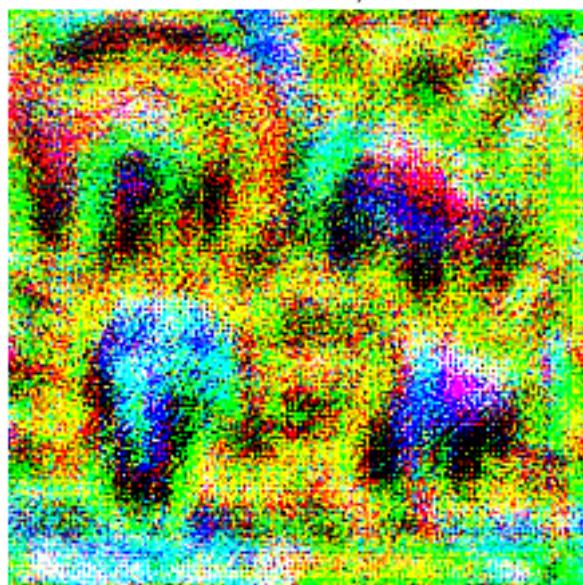
y\_hat:jigsaw puzzle  
y:brown bear, bruin, Ursus arctos  
Iteration 700 / 1000



y\_hat:peacock  
y:brown bear, bruin, Ursus arctos  
Iteration 750 / 1000



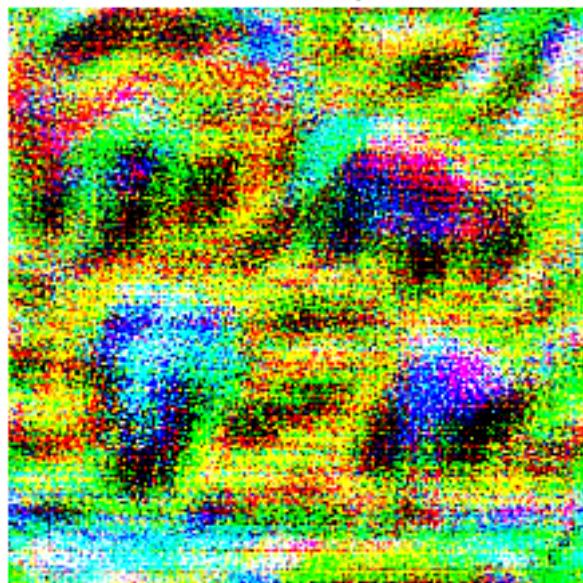
y\_hat:jigsaw puzzle  
y:brown bear, bruin, Ursus arctos  
Iteration 800 / 1000



y\_hat:brown bear, bruin, Ursus arctos  
y:brown bear, bruin, Ursus arctos  
Iteration 850 / 1000



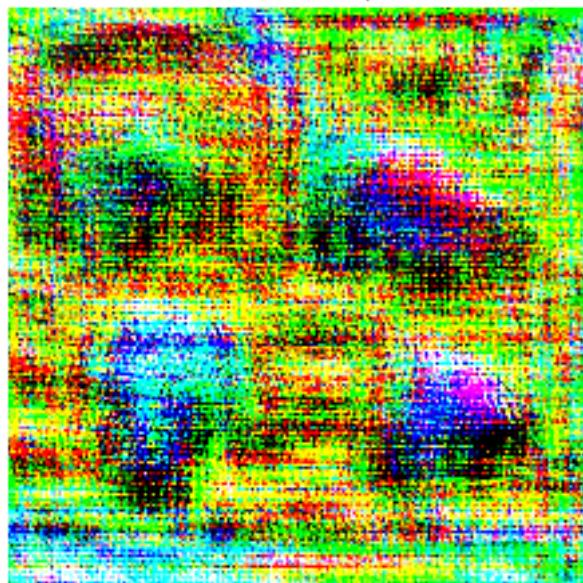
y\_hat:jigsaw puzzle  
y:brown bear, bruin, Ursus arctos  
Iteration 900 / 1000



y\_hat:peacock  
y:brown bear, bruin, Ursus arctos  
Iteration 950 / 1000



y\_hat:jigsaw puzzle  
y:brown bear, bruin, Ursus arctos  
Iteration 1000 / 1000



```
In [22]: # Init du test
    img_ind = 7

    target_y = y[img_ind]

    #target_y = 281 # Tabby cat

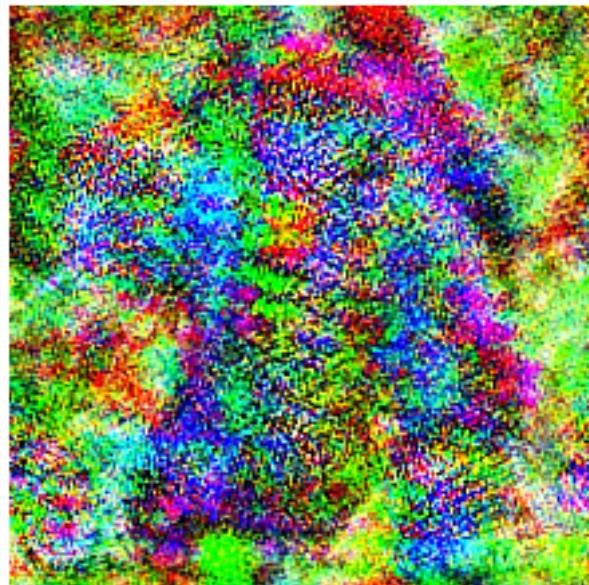
    l2_reg = 1e-2
    learning_rate = 0.1

    X_tensor = torch.Tensor(preprocess(Image.fromarray(X[img_ind])))
    out = create_class_visualization(target_y, model, dtype, init_img=X_tensor, show_every=100)

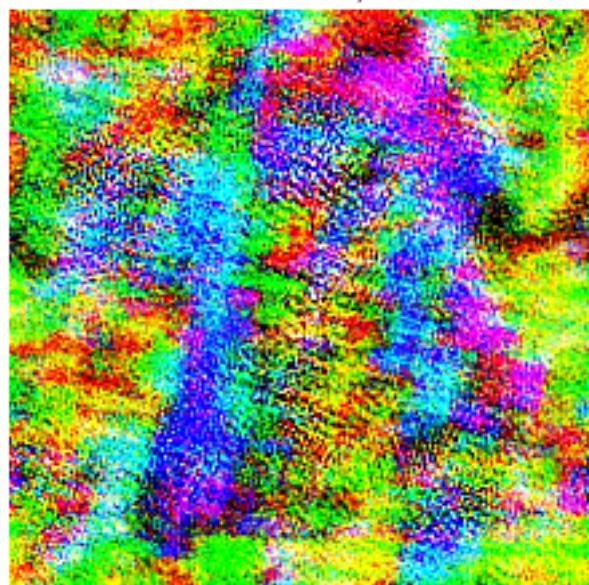
    y_hat:gorilla, Gorilla gorilla
    y:gorilla, Gorilla gorilla
    Iteration 1 / 1000
```



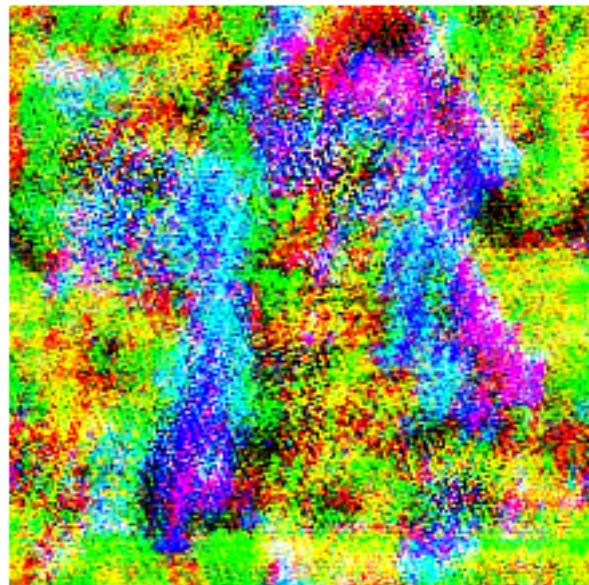
y\_hat:jigsaw puzzle  
y:gorilla, Gorilla gorilla  
Iteration 50 / 1000



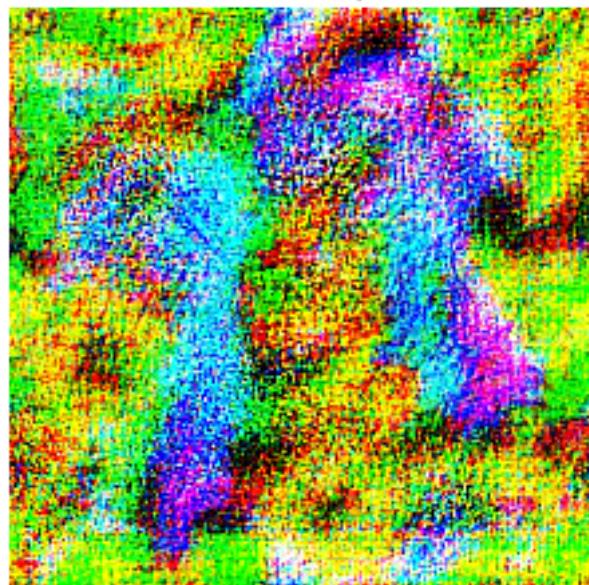
y\_hat:sock  
y:gorilla, Gorilla gorilla  
Iteration 100 / 1000



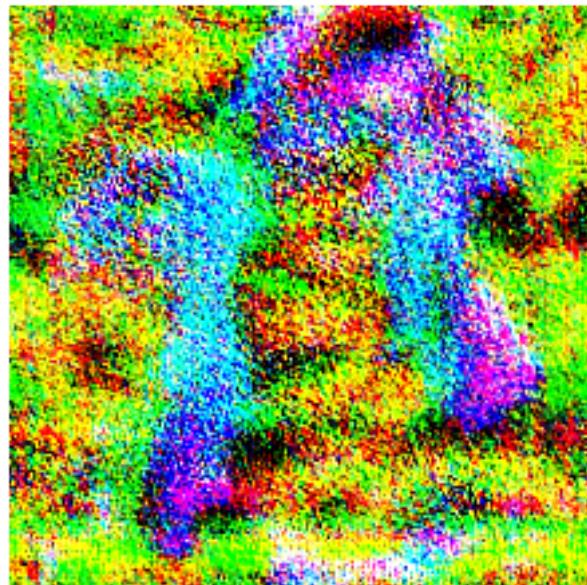
y\_hat:starfish, sea star  
y:gorilla, Gorilla gorilla  
Iteration 150 / 1000



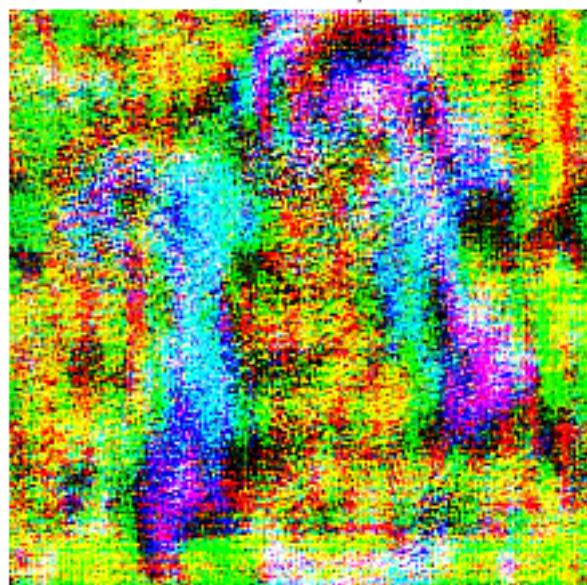
y\_hat:peacock  
y:gorilla, Gorilla gorilla  
Iteration 200 / 1000



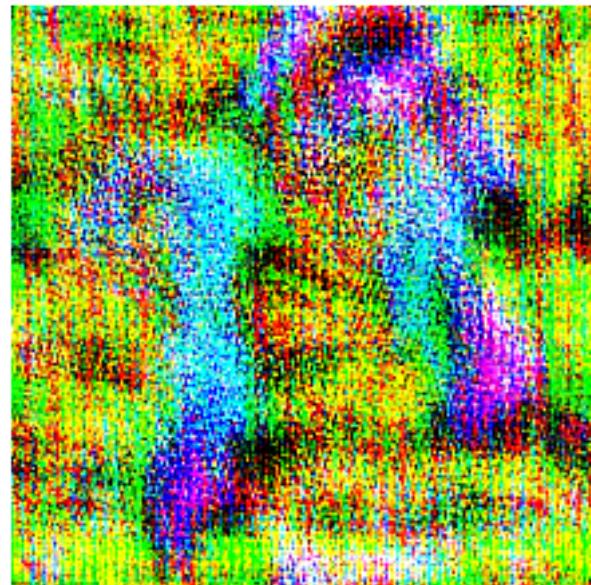
y\_hat:peacock  
y:gorilla, Gorilla gorilla  
Iteration 250 / 1000



y\_hat:sock  
y:gorilla, Gorilla gorilla  
Iteration 300 / 1000



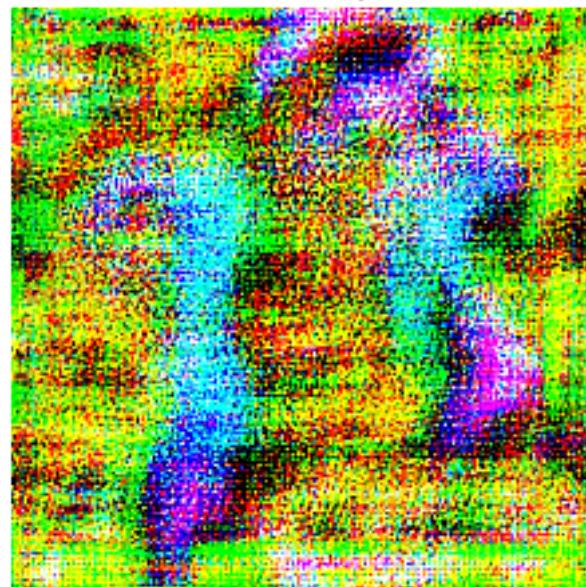
y\_hat:peacock  
y:gorilla, Gorilla gorilla  
Iteration 350 / 1000



y\_hat:peacock  
y:gorilla, Gorilla gorilla  
Iteration 400 / 1000



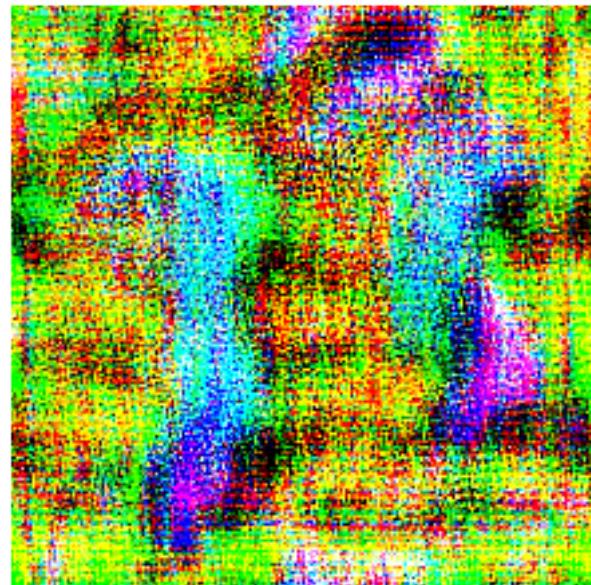
y\_hat:peacock  
y:gorilla, Gorilla gorilla  
Iteration 450 / 1000



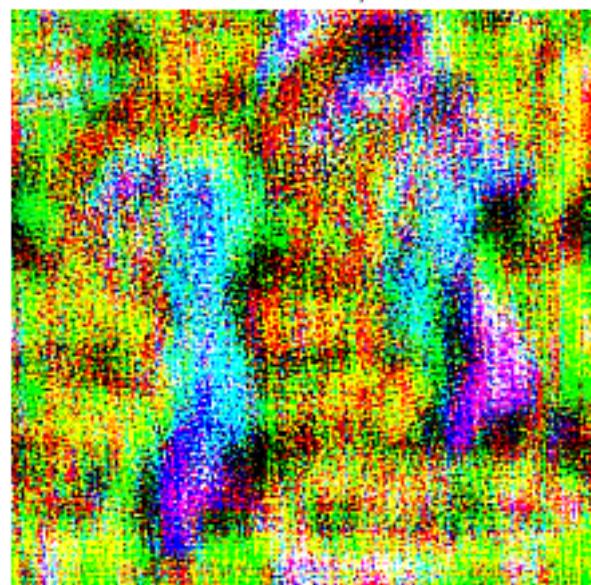
y\_hat:peacock  
y:gorilla, Gorilla gorilla  
Iteration 500 / 1000



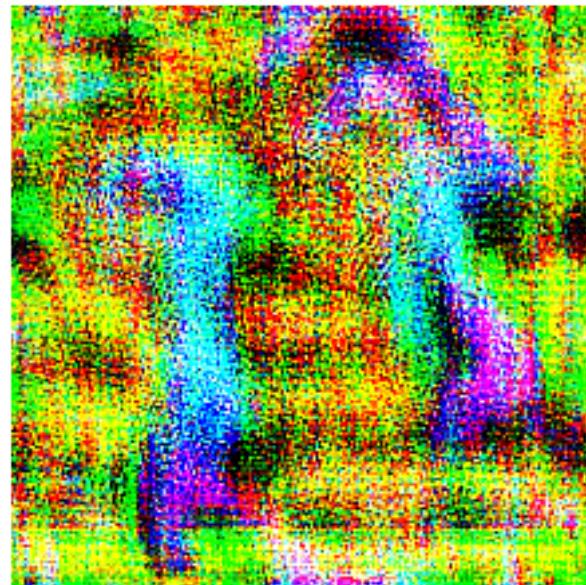
y\_hat:jigsaw puzzle  
y:gorilla, Gorilla gorilla  
Iteration 550 / 1000



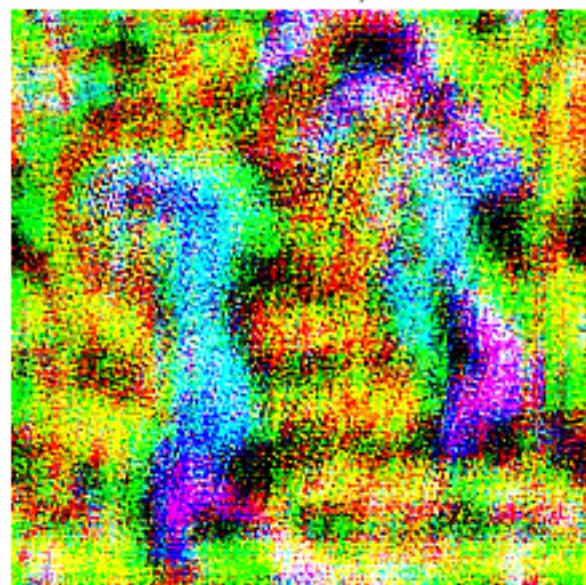
y\_hat:peacock  
y:gorilla, Gorilla gorilla  
Iteration 600 / 1000



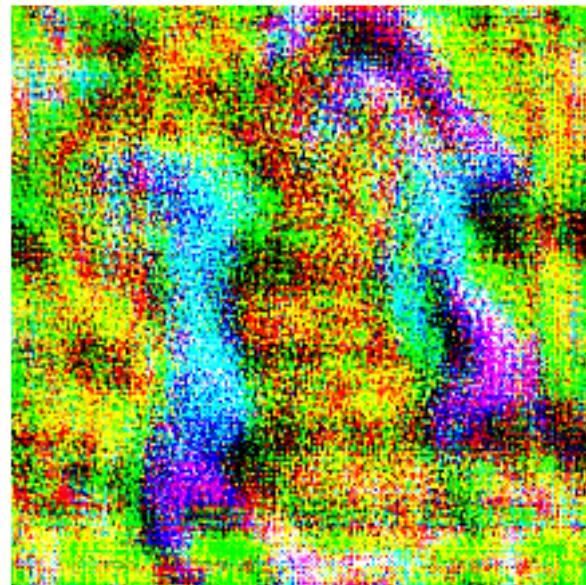
y\_hat:peacock  
y:gorilla, Gorilla gorilla  
Iteration 650 / 1000



y\_hat:peacock  
y:gorilla, Gorilla gorilla  
Iteration 700 / 1000



y\_hat:peacock  
y:gorilla, Gorilla gorilla  
Iteration 750 / 1000



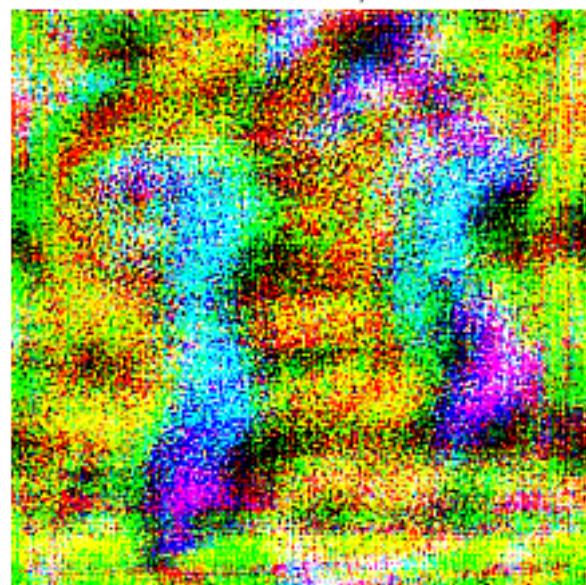
y\_hat:peacock  
y:gorilla, Gorilla gorilla  
Iteration 800 / 1000



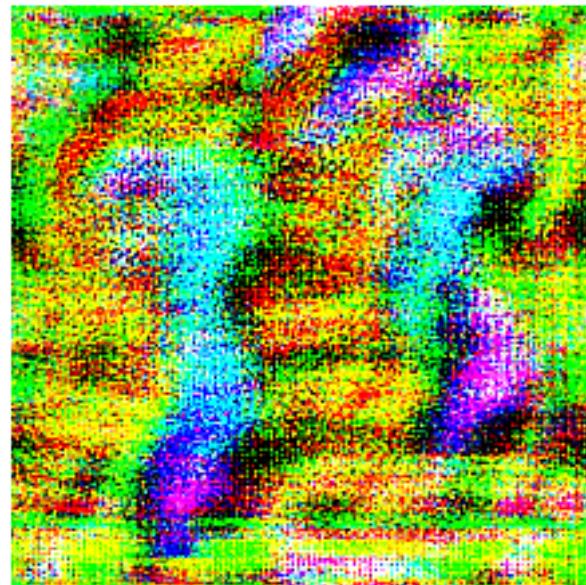
y\_hat:peacock  
y:gorilla, Gorilla gorilla  
Iteration 850 / 1000



y\_hat:coral reef  
y:gorilla, Gorilla gorilla  
Iteration 900 / 1000



y\_hat:peacock  
y:gorilla, Gorilla gorilla  
Iteration 950 / 1000

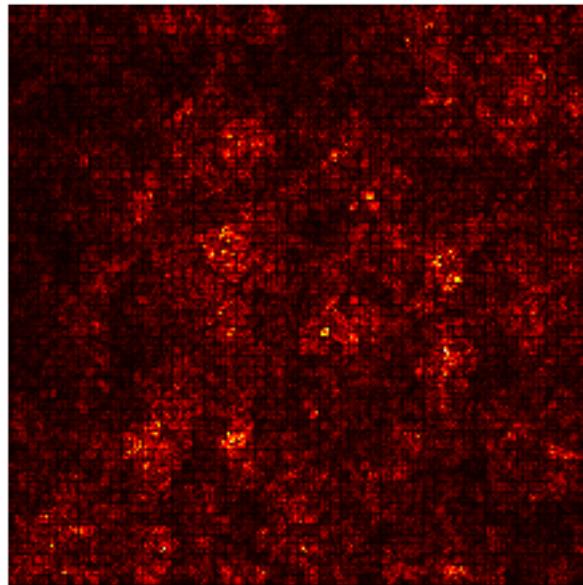


y\_hat:peacock  
y:gorilla, Gorilla gorilla  
Iteration 1000 / 1000



```
In [23]: plt.imshow(compute_saliency_maps(preprocess(out), np.array([target_y]), model)[0] ,  
plt.axis("off")
```

```
Out[23]: (-0.5, 223.5, 223.5, -0.5)
```

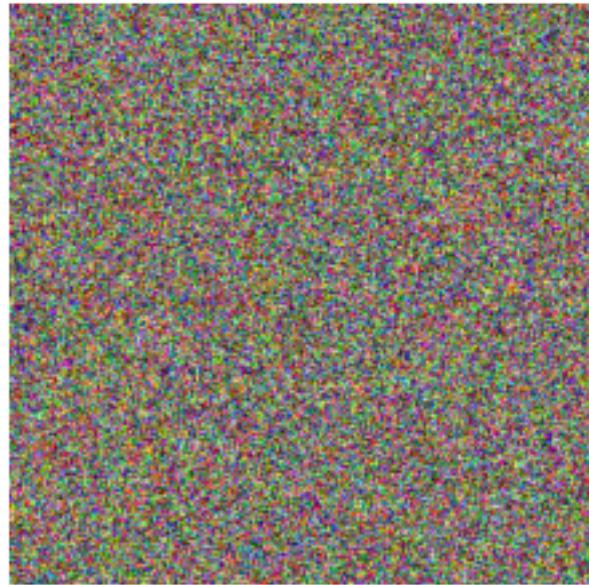


```
In [24]: class_names[target_y]
```

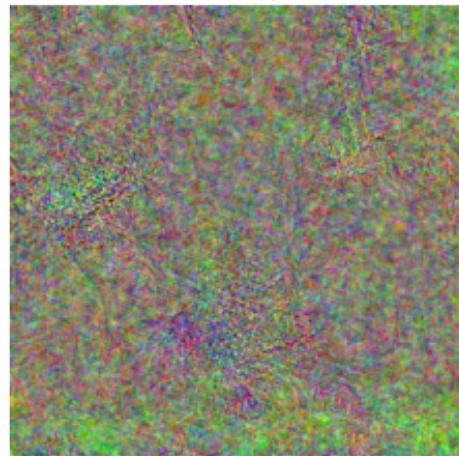
```
Out[24]: 'gorilla, Gorilla gorilla'
```

```
In [25]: dtype = torch.FloatTensor  
# dtype = torch.cuda.FloatTensor # Uncomment this to use GPU  
model.type(dtype)  
  
l2_reg = 1e-3  
learning_rate = 1  
  
num_iterations = 1000  
  
#target_y = 281 # Tabby cat  
# target_y = 187 # Yorkshire Terrier  
#target_y = 76 # Tarantula  
# target_y = 78 # Tick  
# target_y = 683 # Oboe  
target_y = 366 # Gorilla  
# target_y = 604 # Hourglass  
# target_y = np.random.randint(1000) # Classe aléatoire  
out = create_class_visualization(target_y, model, dtype, show_every=50, num_iterations=1000)
```

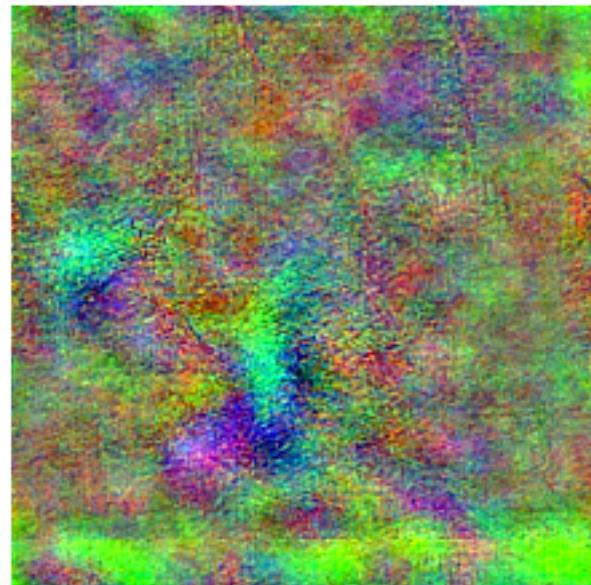
y\_hat:poncho  
y:gorilla, Gorilla gorilla  
Iteration 1 / 1000



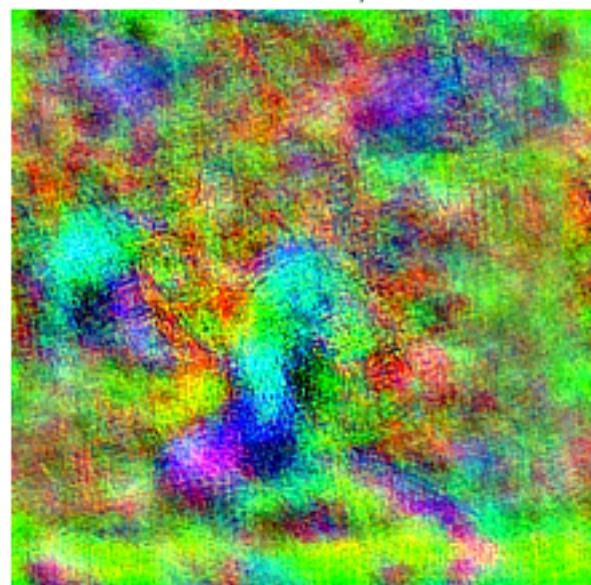
y\_hat:Komodo dragon, Komodo lizard, dragon lizard, giant lizard, Varanus komodoensis  
y:gorilla, Gorilla gorilla  
Iteration 50 / 1000



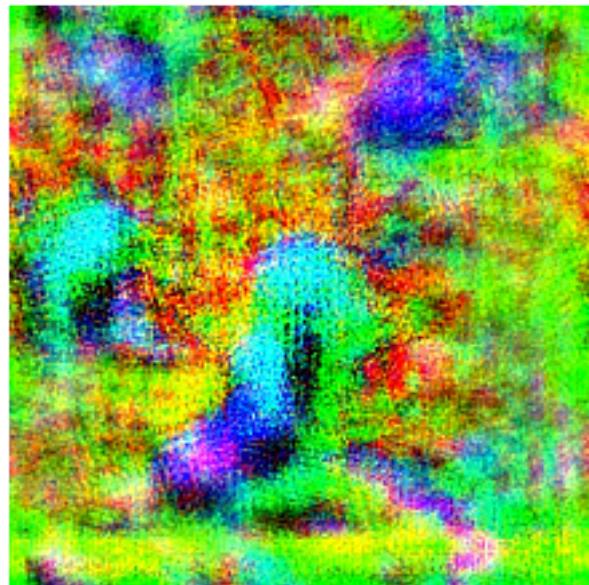
y\_hat:peacock  
y:gorilla, Gorilla gorilla  
Iteration 100 / 1000



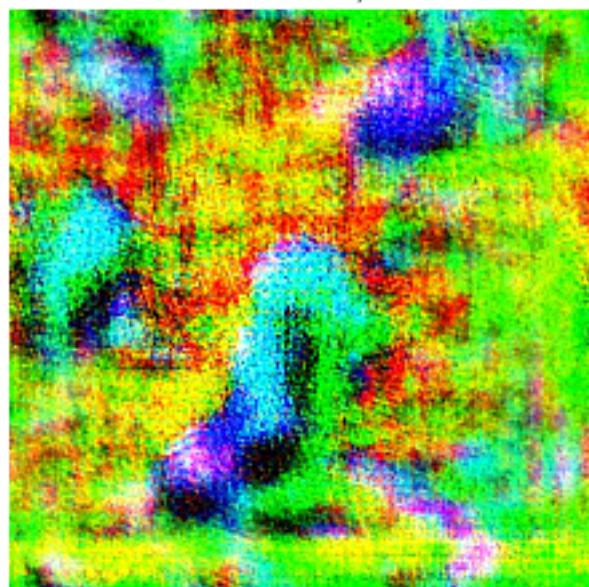
y\_hat:peacock  
y:gorilla, Gorilla gorilla  
Iteration 150 / 1000



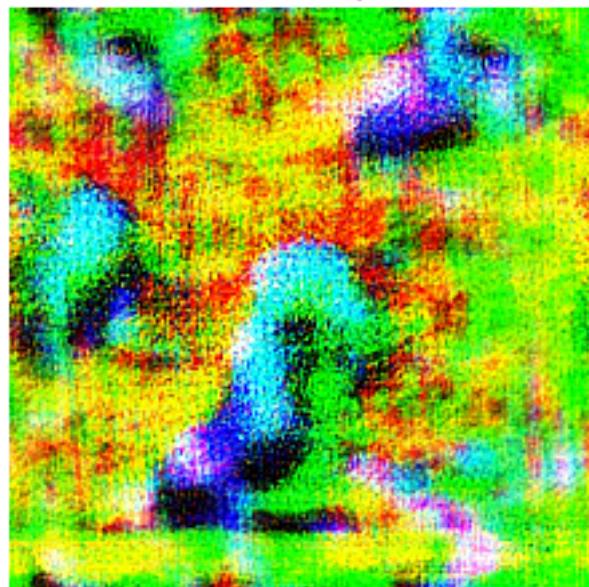
y\_hat:peacock  
y:gorilla, Gorilla gorilla  
Iteration 200 / 1000



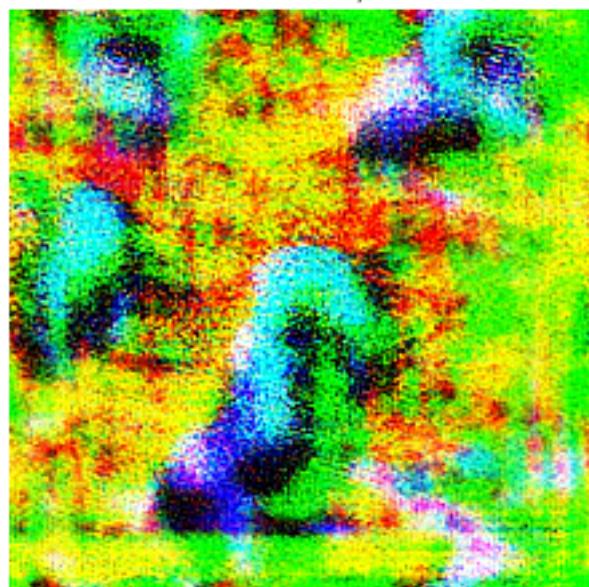
y\_hat:peacock  
y:gorilla, Gorilla gorilla  
Iteration 250 / 1000



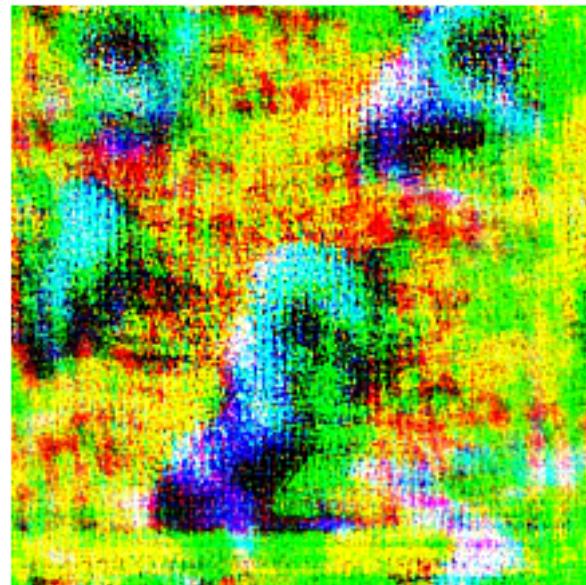
y\_hat:peacock  
y:gorilla, Gorilla gorilla  
Iteration 300 / 1000



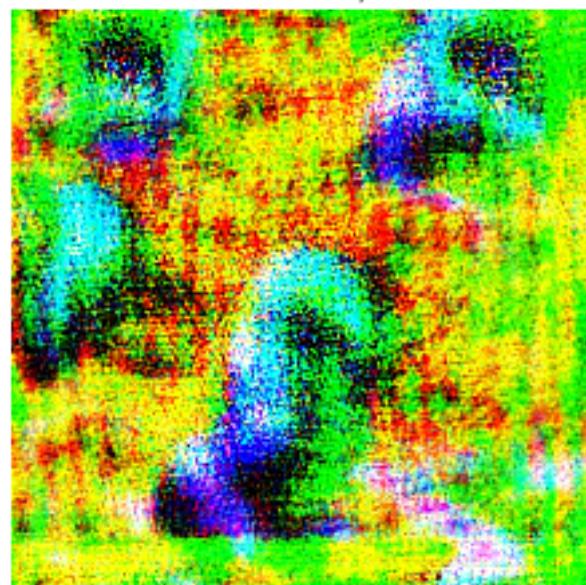
y\_hat:gorilla, Gorilla gorilla  
y:gorilla, Gorilla gorilla  
Iteration 350 / 1000



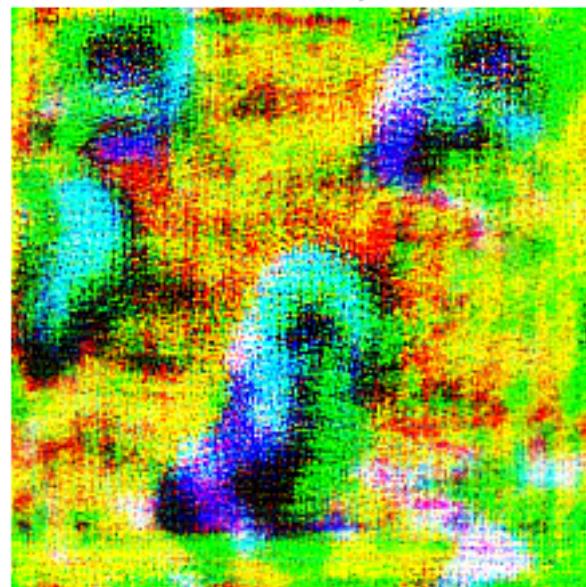
y\_hat:peacock  
y:gorilla, Gorilla gorilla  
Iteration 400 / 1000



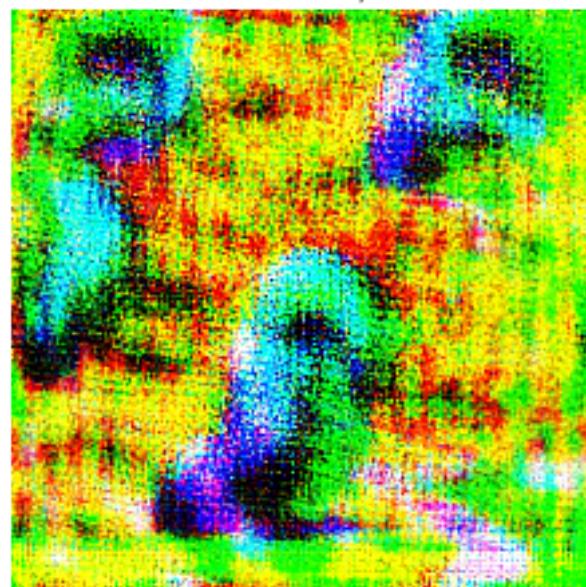
y\_hat:peacock  
y:gorilla, Gorilla gorilla  
Iteration 450 / 1000



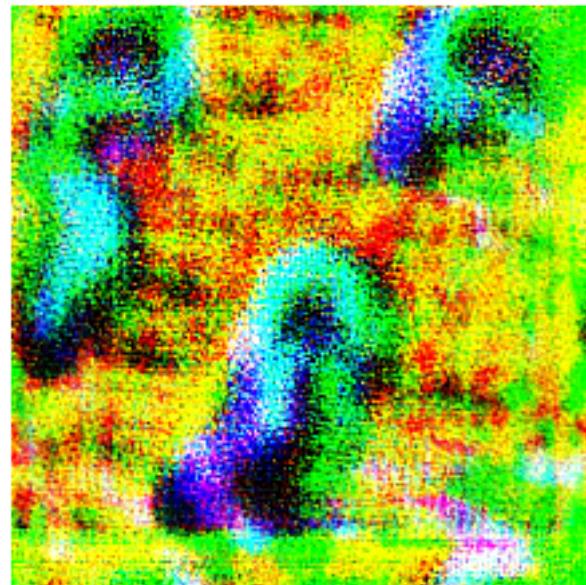
y\_hat:peacock  
y:gorilla, Gorilla gorilla  
Iteration 500 / 1000



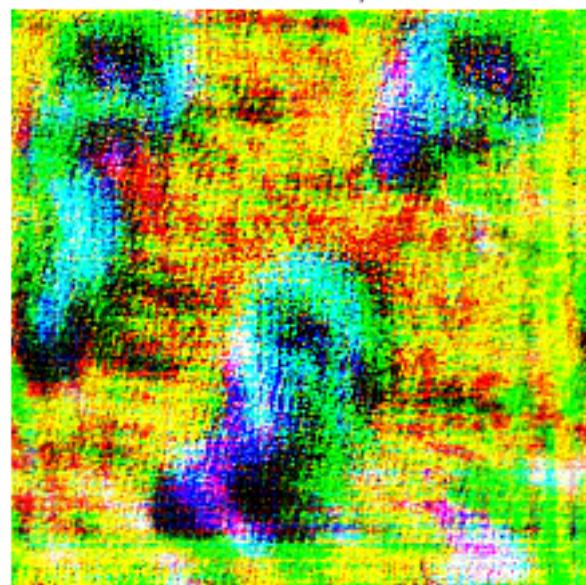
y\_hat:peacock  
y:gorilla, Gorilla gorilla  
Iteration 550 / 1000



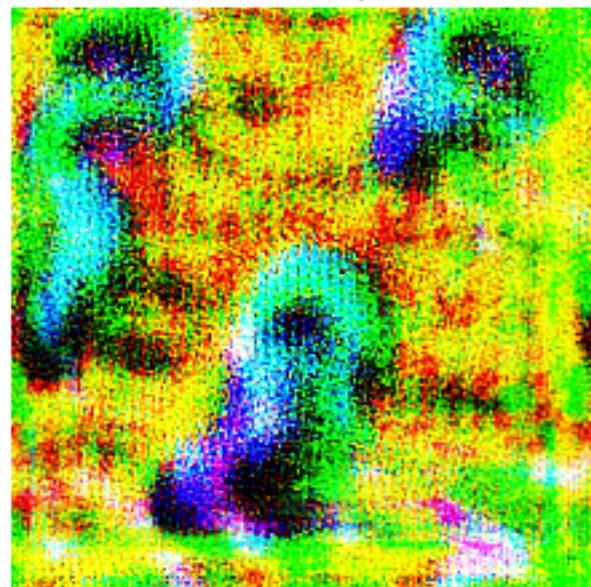
y\_hat:peacock  
y:gorilla, Gorilla gorilla  
Iteration 600 / 1000



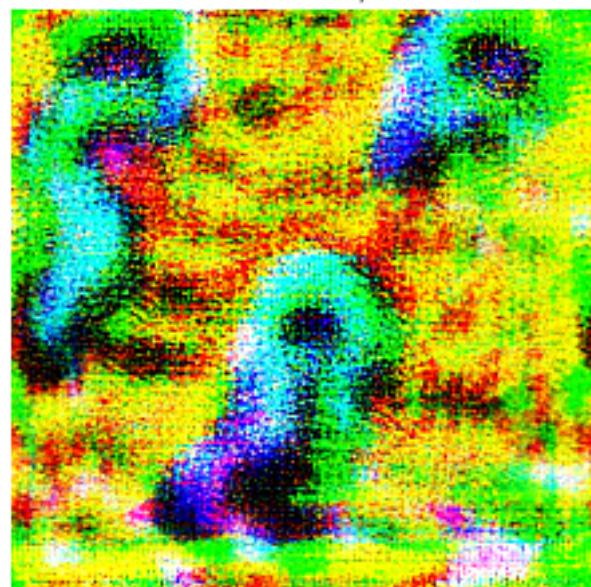
y\_hat:peacock  
y:gorilla, Gorilla gorilla  
Iteration 650 / 1000



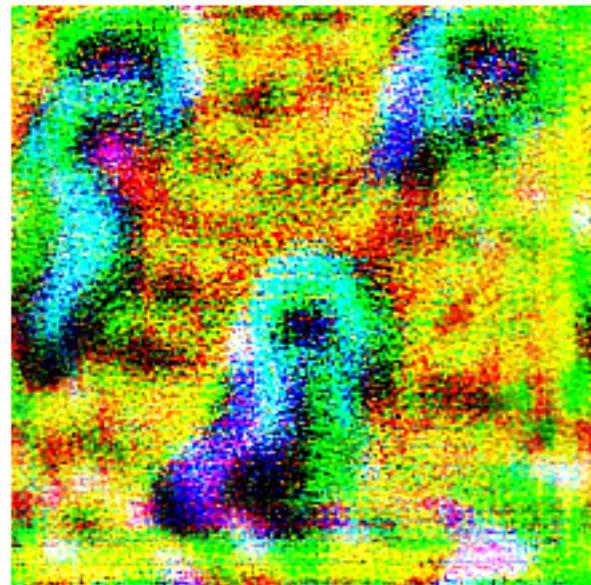
y\_hat:peacock  
y:gorilla, Gorilla gorilla  
Iteration 700 / 1000



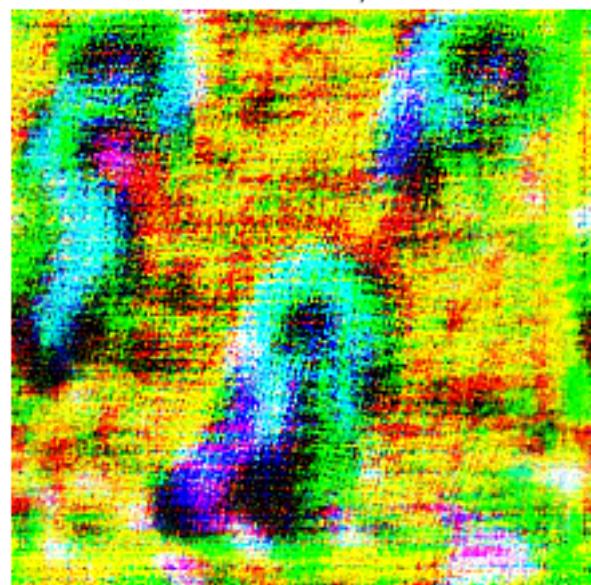
y\_hat:peacock  
y:gorilla, Gorilla gorilla  
Iteration 750 / 1000



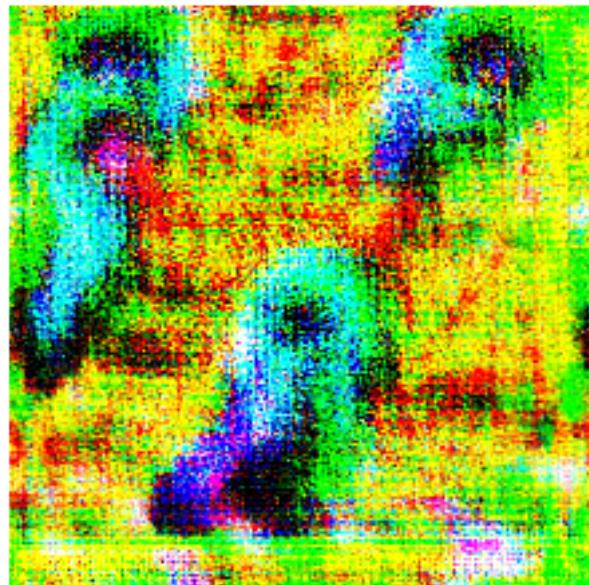
y\_hat:peacock  
y:gorilla, Gorilla gorilla  
Iteration 800 / 1000



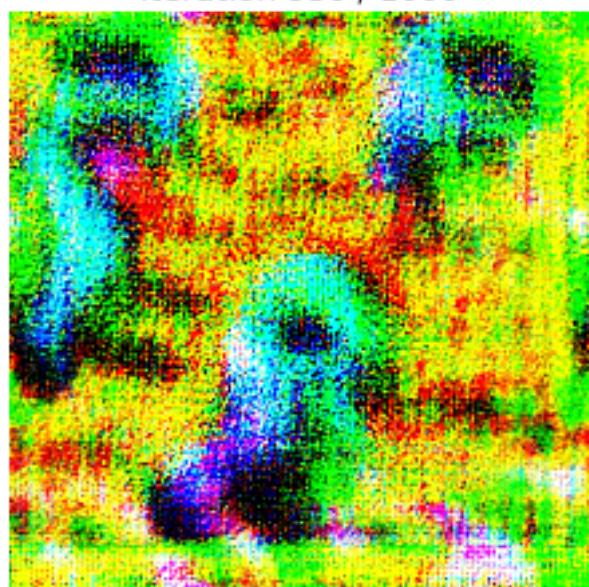
y\_hat:peacock  
y:gorilla, Gorilla gorilla  
Iteration 850 / 1000

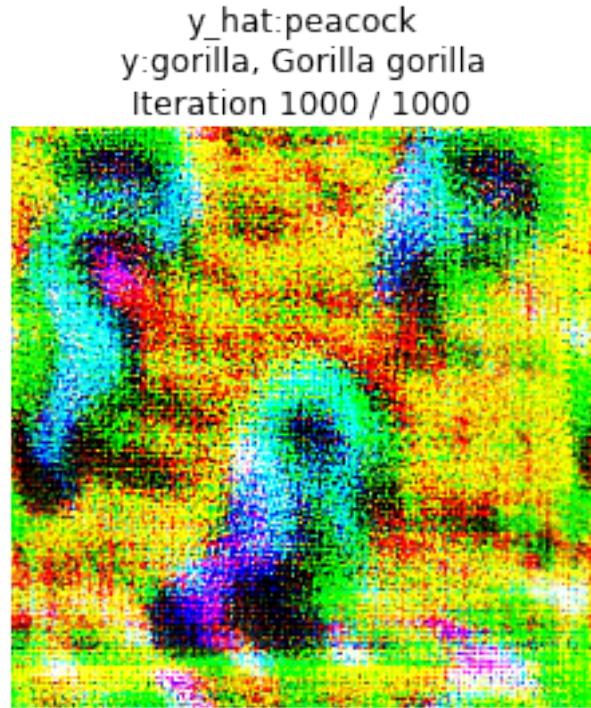


y\_hat:jigsaw puzzle  
y:gorilla, Gorilla gorilla  
Iteration 900 / 1000



y\_hat:peacock  
y:gorilla, Gorilla gorilla  
Iteration 950 / 1000





```
In [26]: dtype = torch.FloatTensor
# dtype = torch.cuda.FloatTensor # Uncomment this to use GPU
model.type(dtype)

l2_reg = 1e-3
learning_rate = 5

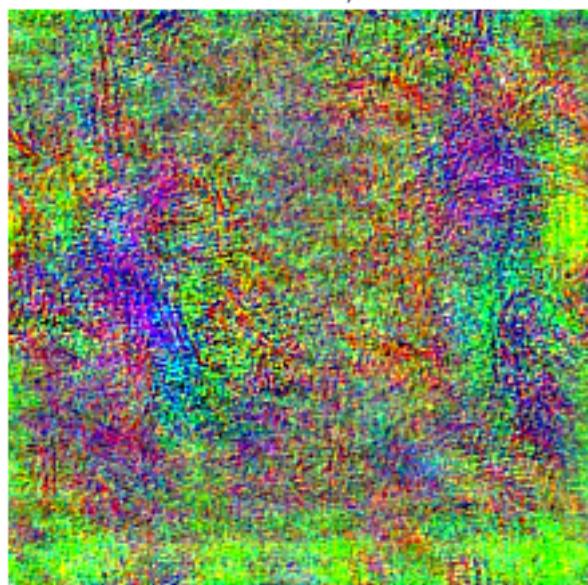
num_iterations = 1000

#target_y = 281 # Tabby cat
# target_y = 187 # Yorkshire Terrier
#target_y = 76 # Tarantula
# target_y = 78 # Tick
# target_y = 683 # Oboe
target_y = 366 # Gorilla
# target_y = 604 # Hourglass
# target_y = np.random.randint(1000) # Classe aléatoire
out = create_class_visualization(target_y, model, dtype, show_every=50, num_iterations)
```

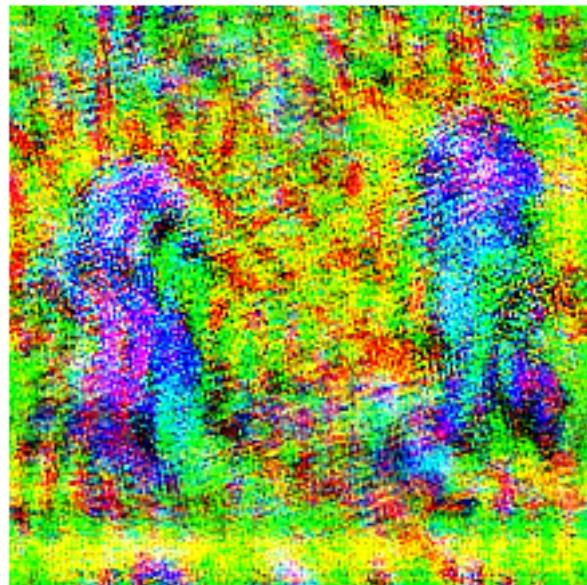
y\_hat:doormat, welcome mat  
y:gorilla, Gorilla gorilla  
Iteration 1 / 1000



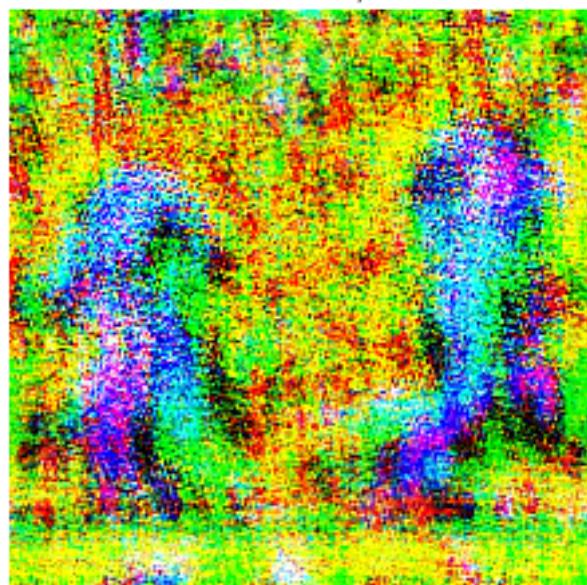
y\_hat:peacock  
y:gorilla, Gorilla gorilla  
Iteration 50 / 1000



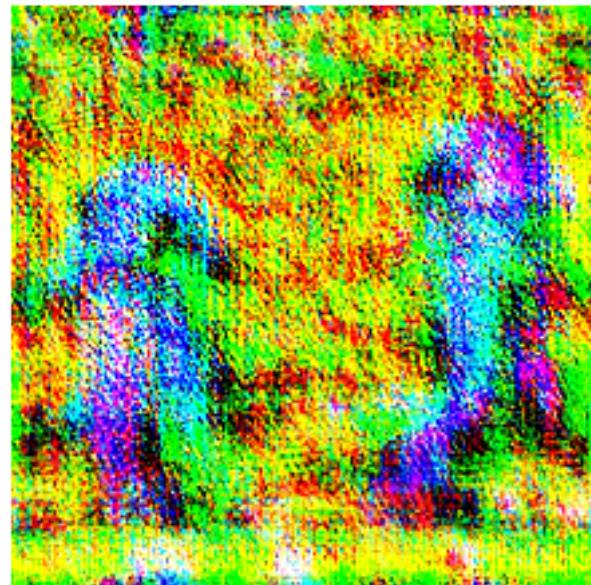
y\_hat:peacock  
y:gorilla, Gorilla gorilla  
Iteration 100 / 1000



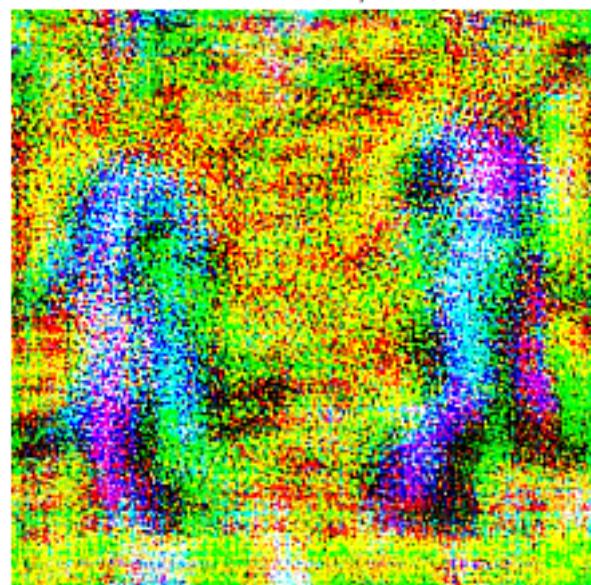
y\_hat:peacock  
y:gorilla, Gorilla gorilla  
Iteration 150 / 1000



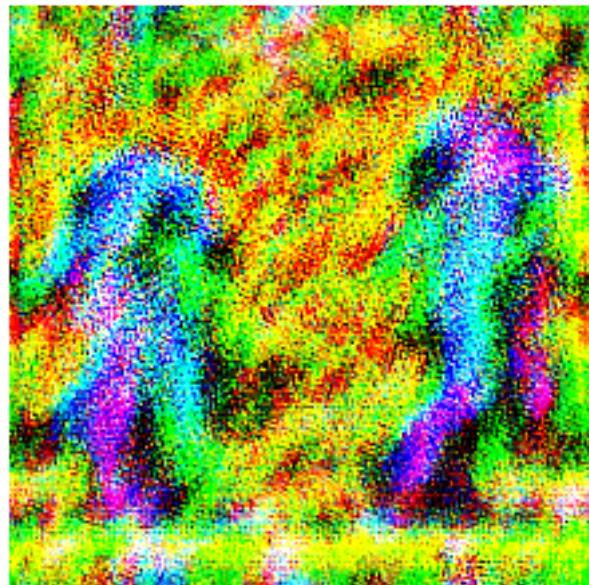
y\_hat:peacock  
y:gorilla, Gorilla gorilla  
Iteration 200 / 1000



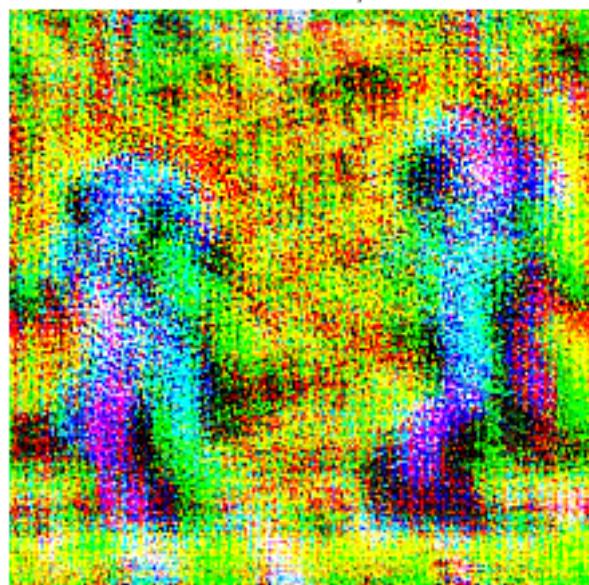
y\_hat:peacock  
y:gorilla, Gorilla gorilla  
Iteration 250 / 1000



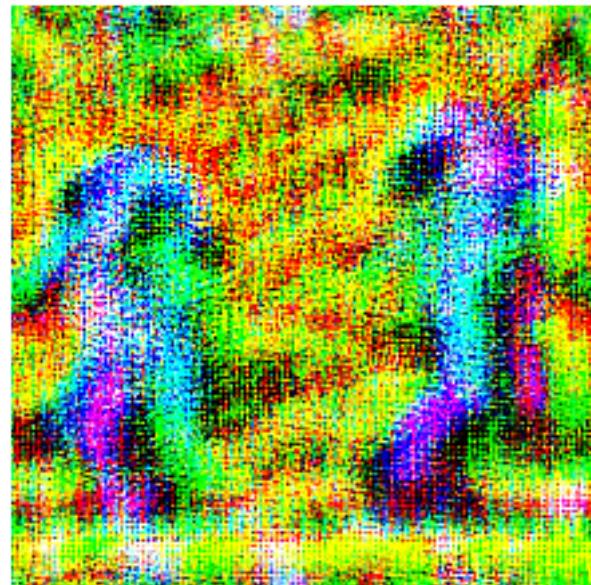
y\_hat:peacock  
y:gorilla, Gorilla gorilla  
Iteration 300 / 1000



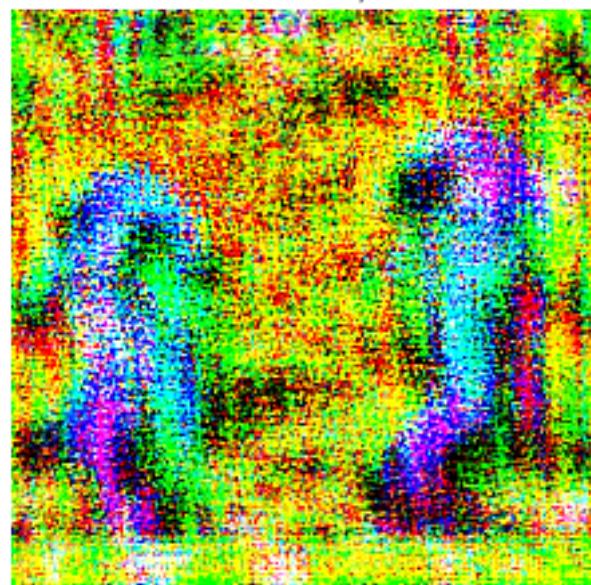
y\_hat:peacock  
y:gorilla, Gorilla gorilla  
Iteration 350 / 1000



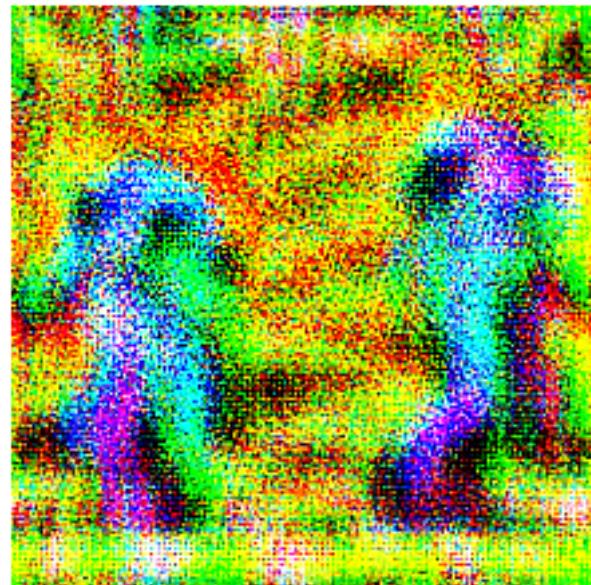
y\_hat:peacock  
y:gorilla, Gorilla gorilla  
Iteration 400 / 1000



y\_hat:peacock  
y:gorilla, Gorilla gorilla  
Iteration 450 / 1000



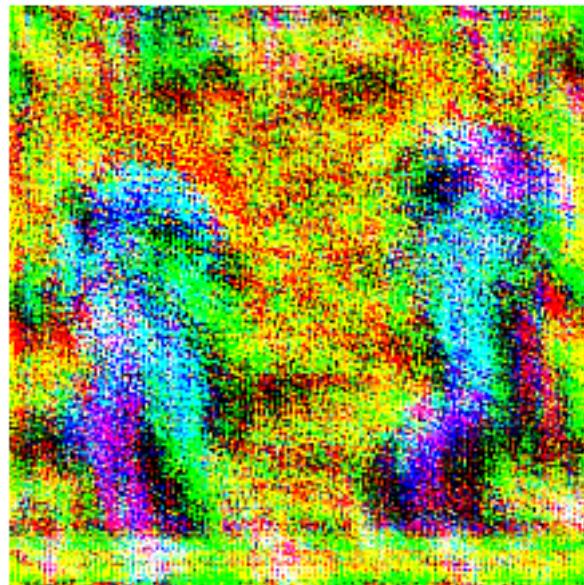
y\_hat:peacock  
y:gorilla, Gorilla gorilla  
Iteration 500 / 1000



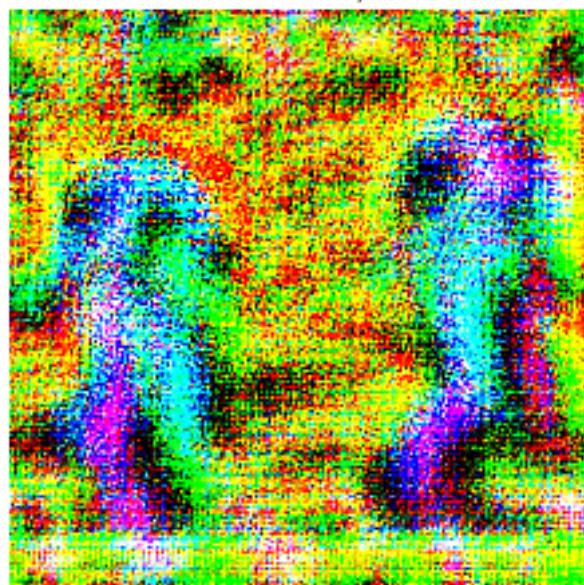
y\_hat:peacock  
y:gorilla, Gorilla gorilla  
Iteration 550 / 1000



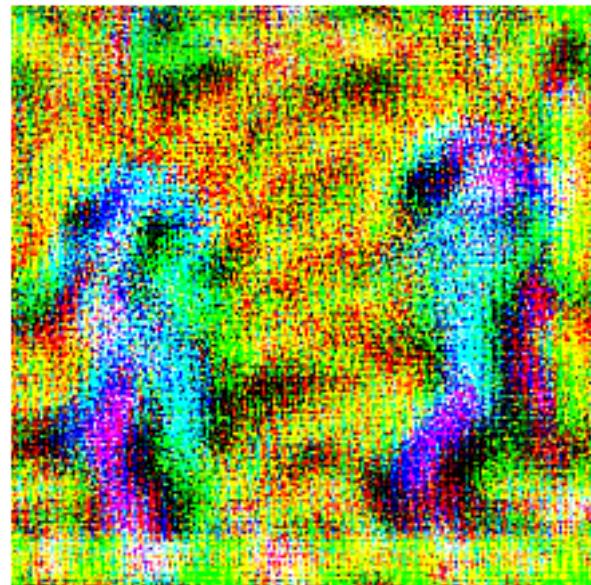
y\_hat:peacock  
y:gorilla, Gorilla gorilla  
Iteration 600 / 1000



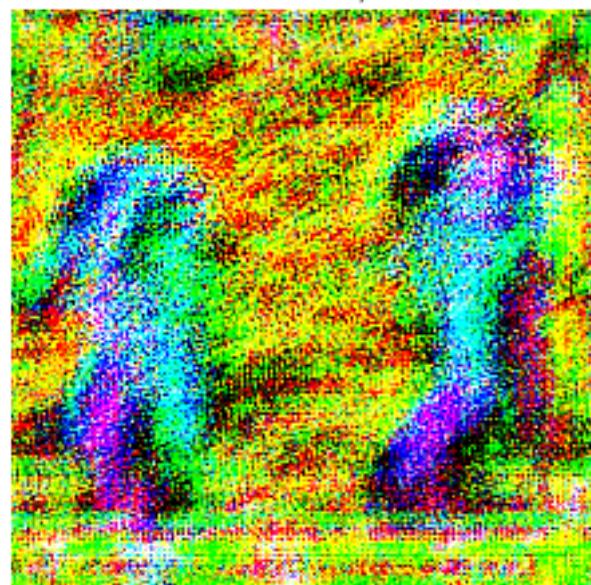
y\_hat:standard poodle  
y:gorilla, Gorilla gorilla  
Iteration 650 / 1000



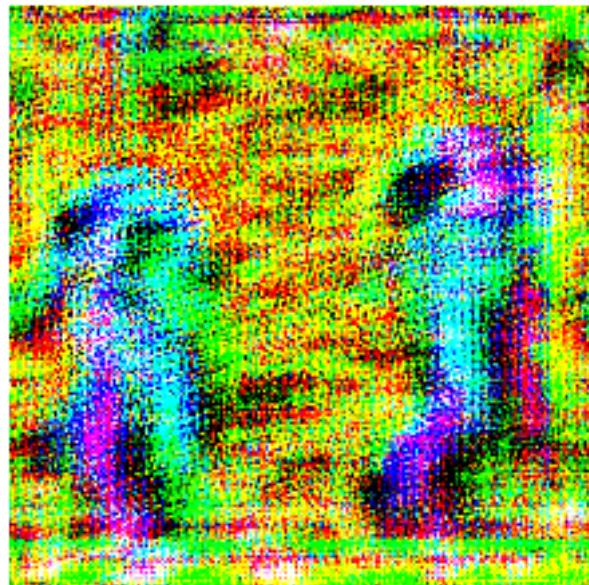
y\_hat:peacock  
y:gorilla, Gorilla gorilla  
Iteration 700 / 1000



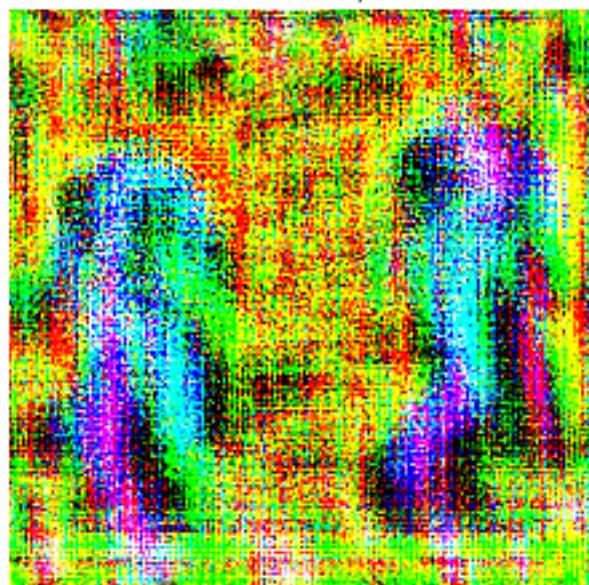
y\_hat:jigsaw puzzle  
y:gorilla, Gorilla gorilla  
Iteration 750 / 1000



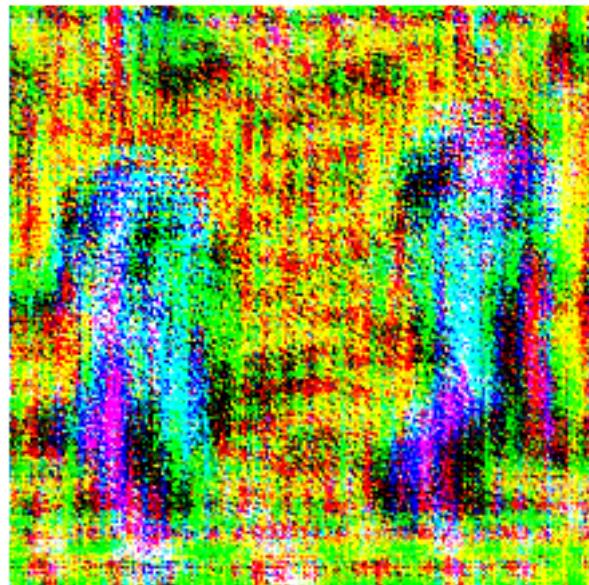
y\_hat:peacock  
y:gorilla, Gorilla gorilla  
Iteration 800 / 1000



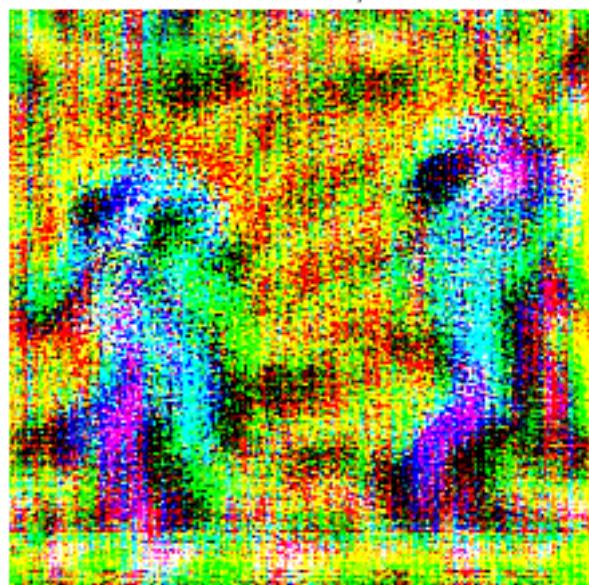
y\_hat:peacock  
y:gorilla, Gorilla gorilla  
Iteration 850 / 1000



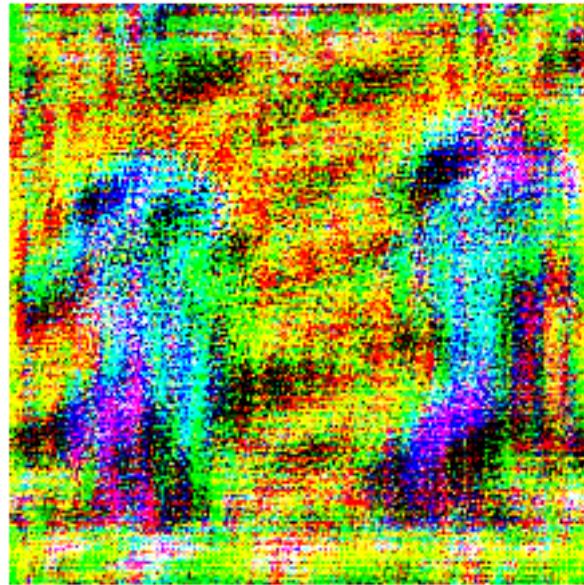
y\_hat:standard poodle  
y:gorilla, Gorilla gorilla  
Iteration 900 / 1000



y\_hat:prayer rug, prayer mat  
y:gorilla, Gorilla gorilla  
Iteration 950 / 1000

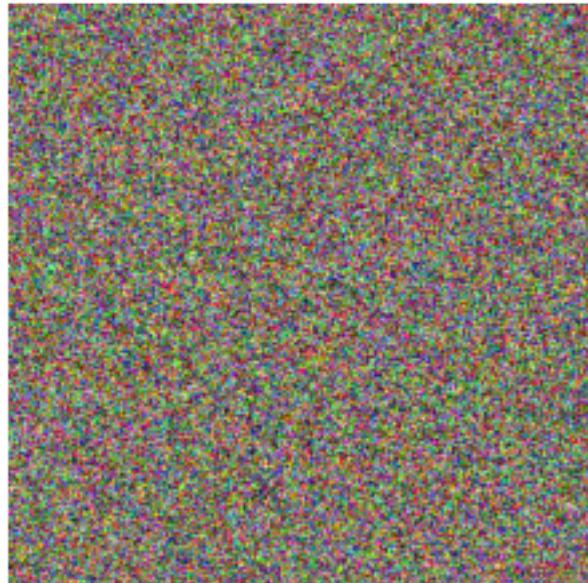


```
y_hat:prayer rug, prayer mat  
y:gorilla, Gorilla gorilla  
Iteration 1000 / 1000
```



```
In [27]: dtype = torch.FloatTensor  
# dtype = torch.cuda.FloatTensor # Uncomment this to use GPU  
model.type(dtype)  
  
l2_reg = 1e-3  
learning_rate = 0.01  
  
num_iterations = 1000  
  
#target_y = 281 # Tabby cat  
# target_y = 187 # Yorkshire Terrier  
#target_y = 76 # Tarantula  
# target_y = 78 # Tick  
# target_y = 683 # Oboe  
target_y = 366 # Gorilla  
# target_y = 604 # Hourglass  
# target_y = np.random.randint(1000) # Classe aléatoire  
out = create_class_visualization(target_y, model, dtype, show_every=50, num_iterations)
```

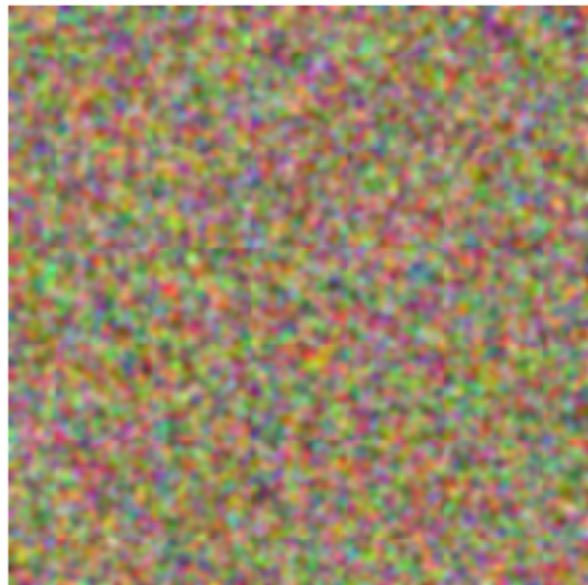
y\_hat:doormat, welcome mat  
y:gorilla, Gorilla gorilla  
Iteration 1 / 1000



y\_hat:doormat, welcome mat  
y:gorilla, Gorilla gorilla  
Iteration 50 / 1000



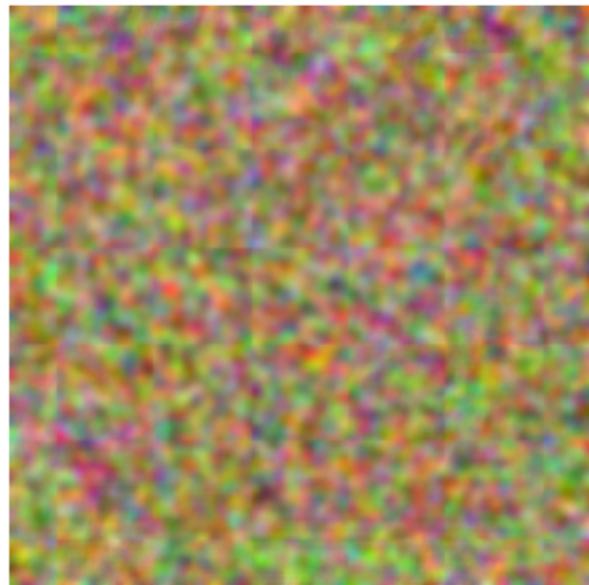
y\_hat:doormat, welcome mat  
y:gorilla, Gorilla gorilla  
Iteration 100 / 1000



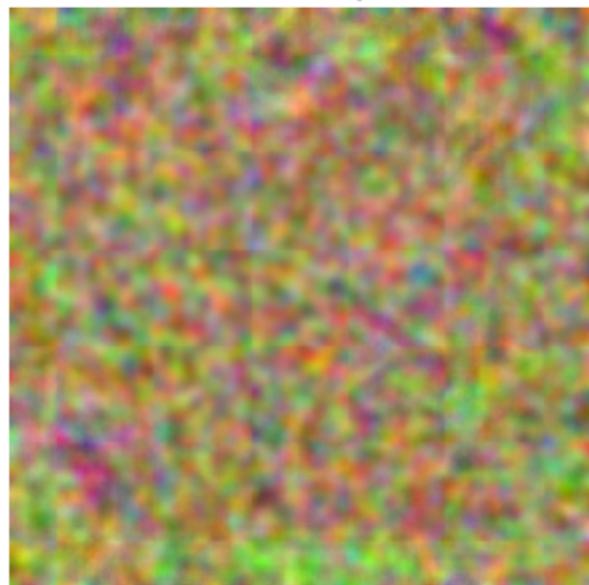
y\_hat:croquet ball  
y:gorilla, Gorilla gorilla  
Iteration 150 / 1000



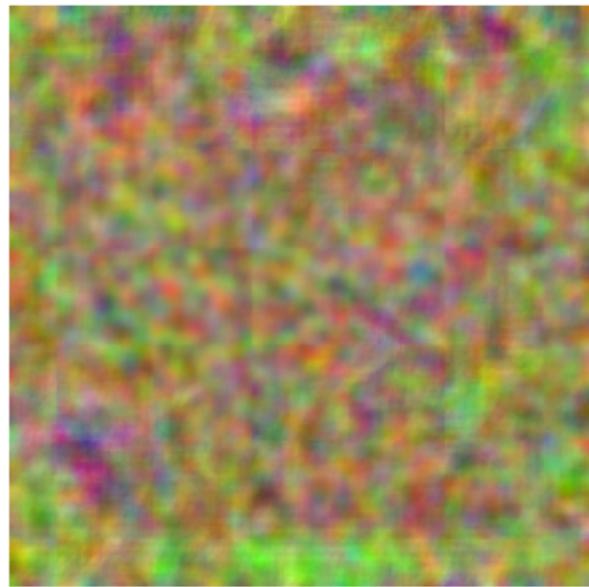
y\_hat:hare  
y:gorilla, Gorilla gorilla  
Iteration 200 / 1000



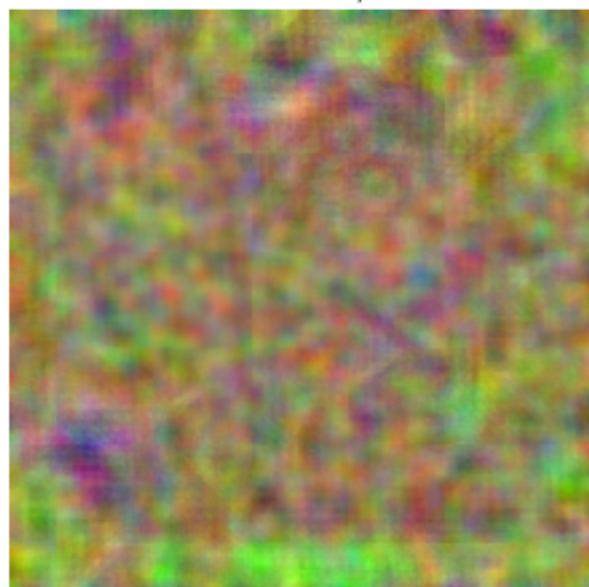
y\_hat:hare  
y:gorilla, Gorilla gorilla  
Iteration 250 / 1000



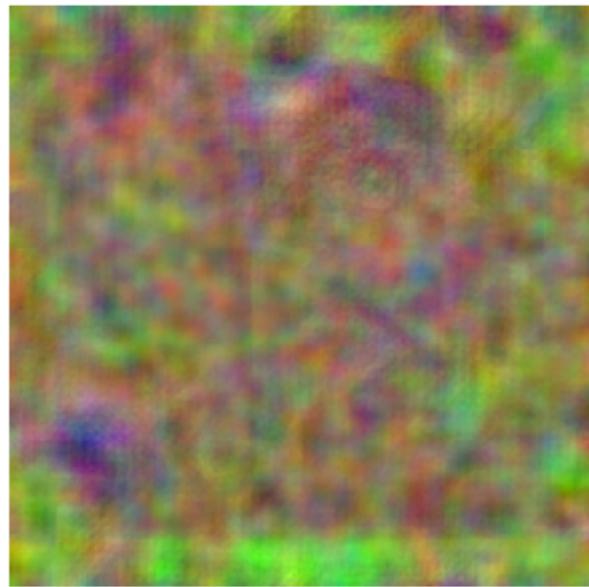
y\_hat:hare  
y:gorilla, Gorilla gorilla  
Iteration 300 / 1000



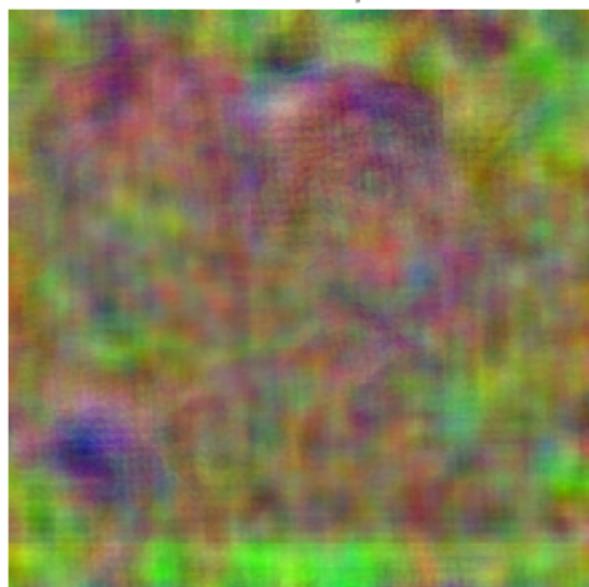
y\_hat:hare  
y:gorilla, Gorilla gorilla  
Iteration 350 / 1000



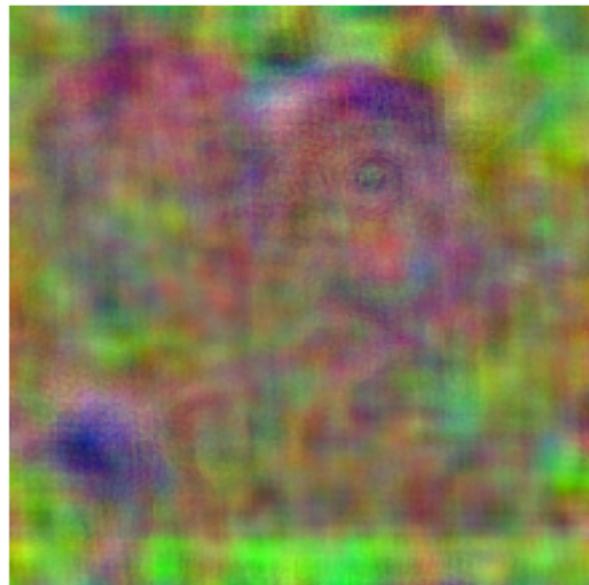
y\_hat: hare  
y:gorilla, Gorilla gorilla  
Iteration 400 / 1000



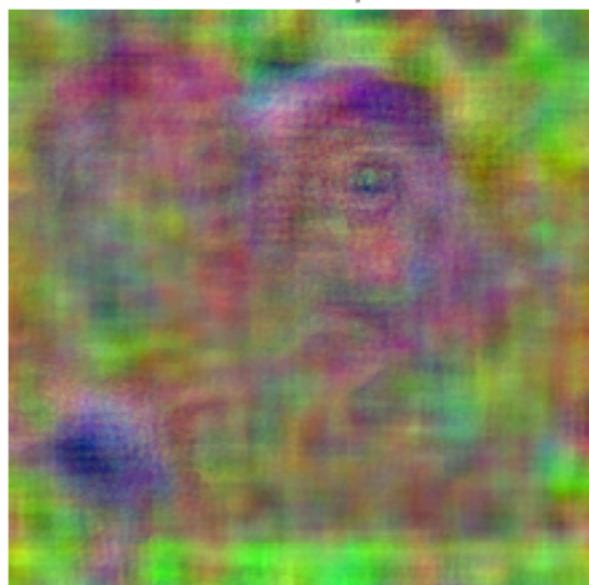
y\_hat: patas, hussar monkey, Erythrocebus patas  
y:gorilla, Gorilla gorilla  
Iteration 450 / 1000



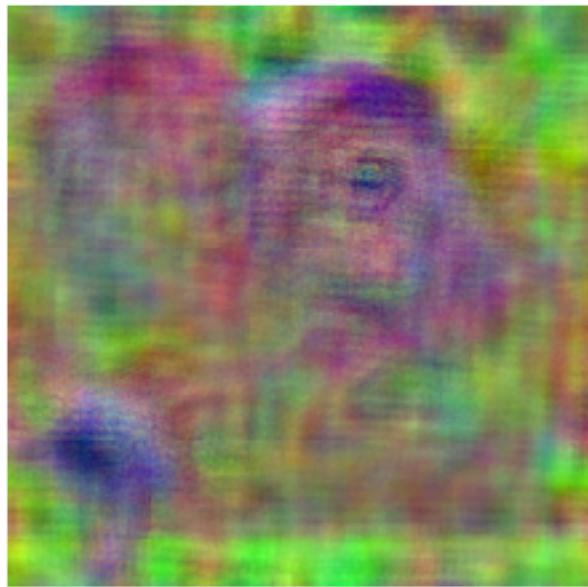
y\_hat:baboon  
y:gorilla, Gorilla gorilla  
Iteration 500 / 1000



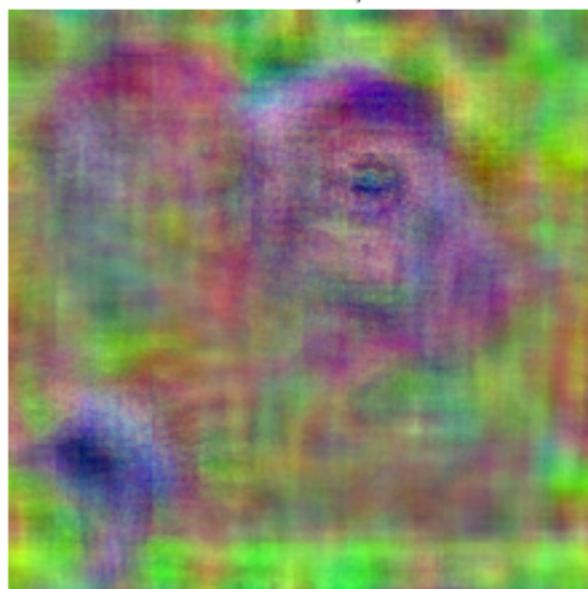
y\_hat:patas, hussar monkey, Erythrocebus patas  
y:gorilla, Gorilla gorilla  
Iteration 550 / 1000



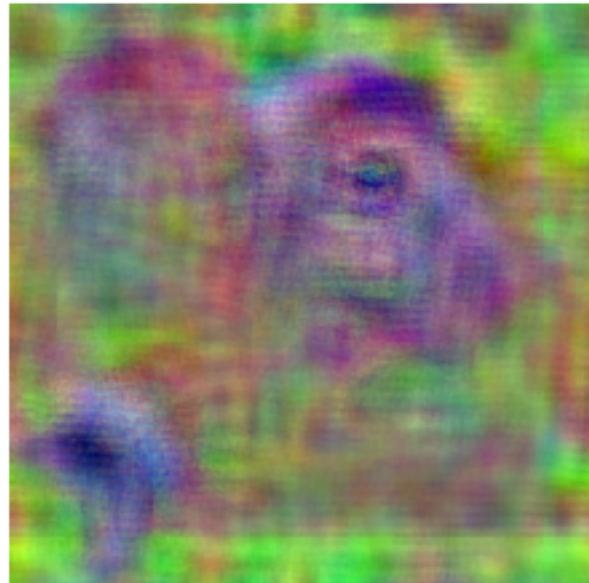
y\_hat: patas, hussar monkey, Erythrocebus patas  
y: gorilla, Gorilla gorilla  
Iteration 600 / 1000



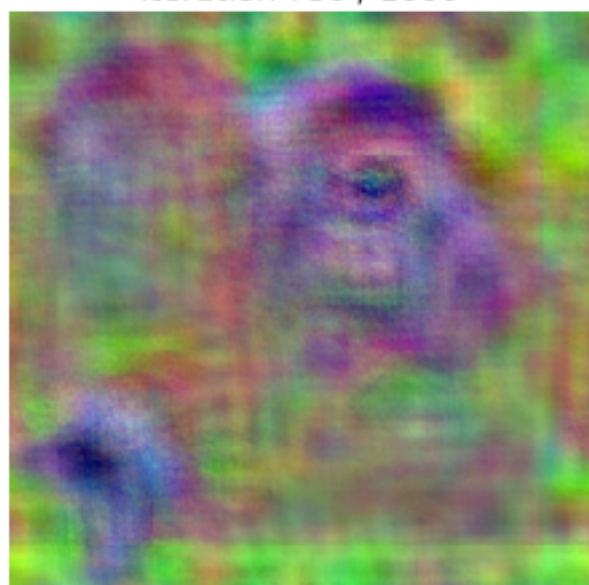
y\_hat: gorilla, Gorilla gorilla  
y: gorilla, Gorilla gorilla  
Iteration 650 / 1000



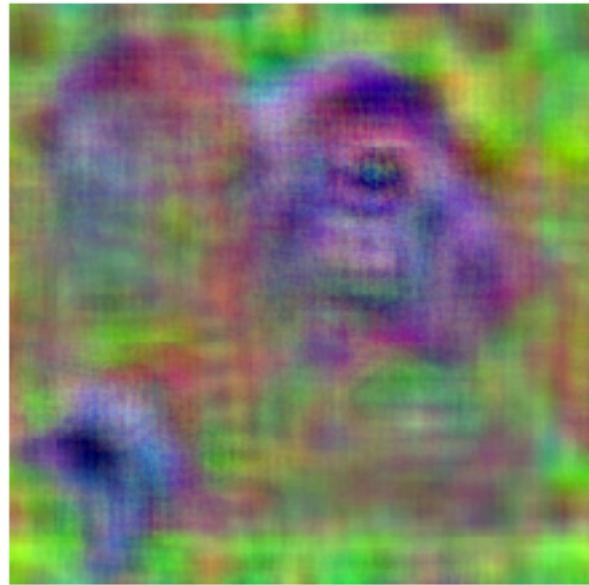
y\_hat:gorilla, Gorilla gorilla  
y:gorilla, Gorilla gorilla  
Iteration 700 / 1000



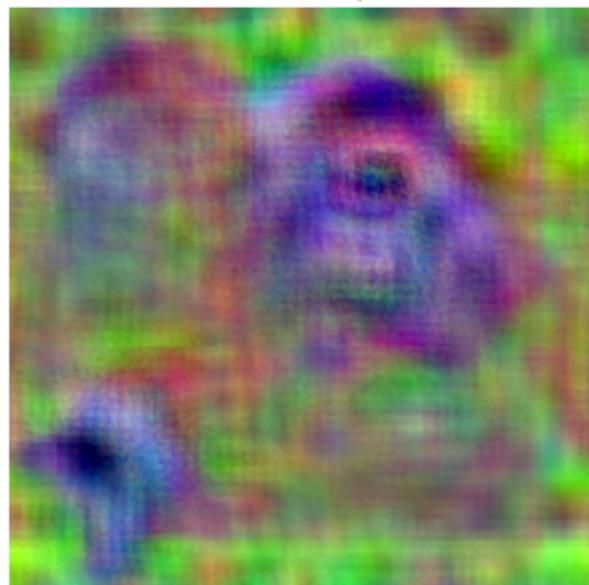
y\_hat:macaque  
y:gorilla, Gorilla gorilla  
Iteration 750 / 1000



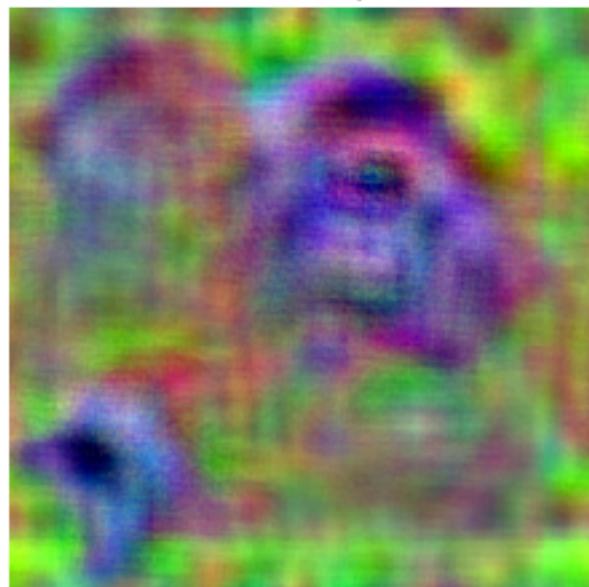
y\_hat:gorilla, Gorilla gorilla  
y:gorilla, Gorilla gorilla  
Iteration 800 / 1000



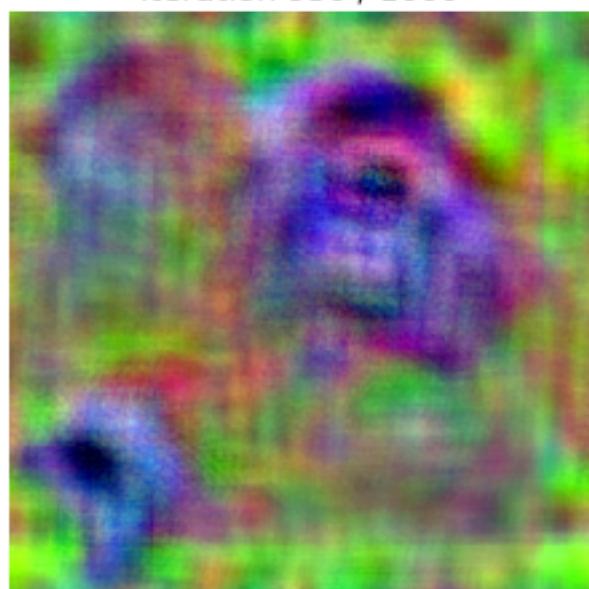
y\_hat:gorilla, Gorilla gorilla  
y:gorilla, Gorilla gorilla  
Iteration 850 / 1000

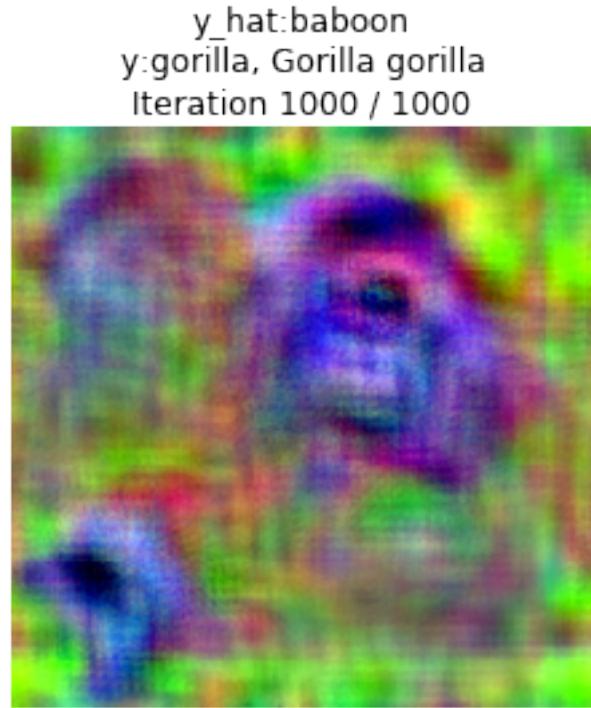


y\_hat:macaque  
y:gorilla, Gorilla gorilla  
Iteration 900 / 1000



y\_hat:gorilla, Gorilla gorilla  
y:gorilla, Gorilla gorilla  
Iteration 950 / 1000





```

In [28]: dtype = torch.FloatTensor
# dtype = torch.cuda.FloatTensor # Uncomment this to use GPU
model.type(dtype)

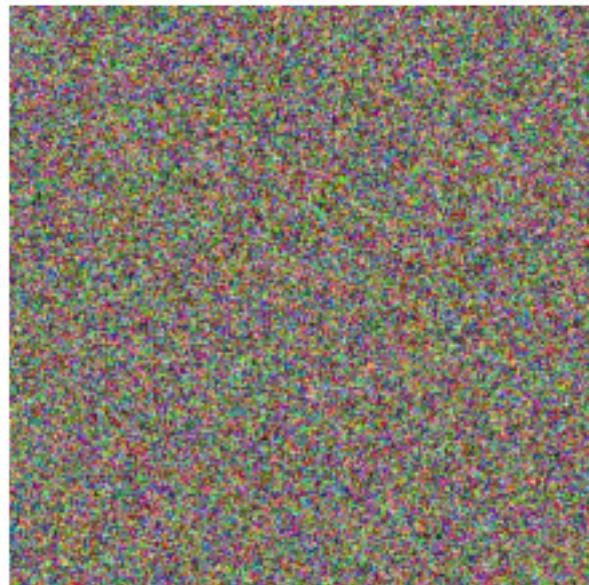
l2_reg = 1e-3
learning_rate = 0.001

num_iterations = 1000

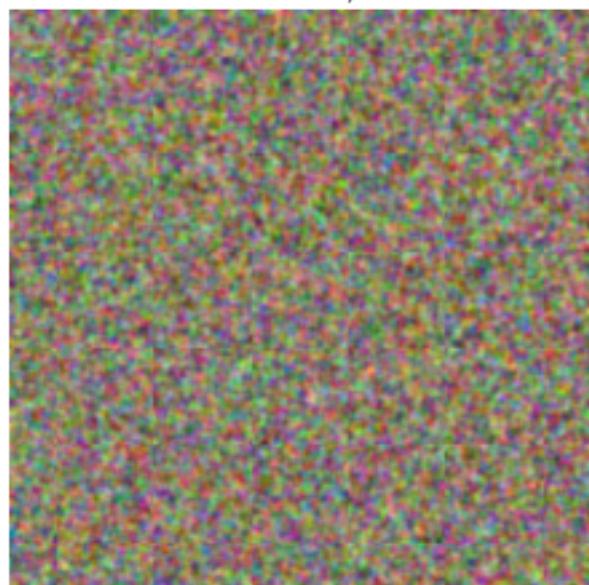
#target_y = 281 # Tabby cat
# target_y = 187 # Yorkshire Terrier
#target_y = 76 # Tarantula
# target_y = 78 # Tick
# target_y = 683 # Oboe
target_y = 366 # Gorilla
# target_y = 604 # Hourglass
# target_y = np.random.randint(1000) # Classe aléatoire
out = create_class_visualization(target_y, model, dtype, show_every=50, num_iterations)

```

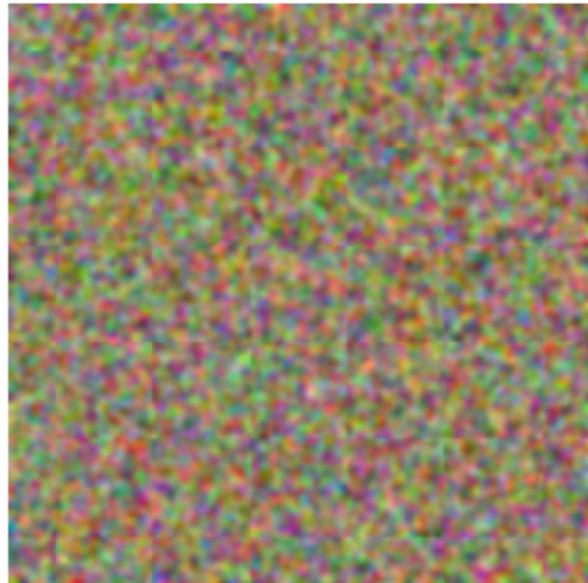
y\_hat:poncho  
y:gorilla, Gorilla gorilla  
Iteration 1 / 1000



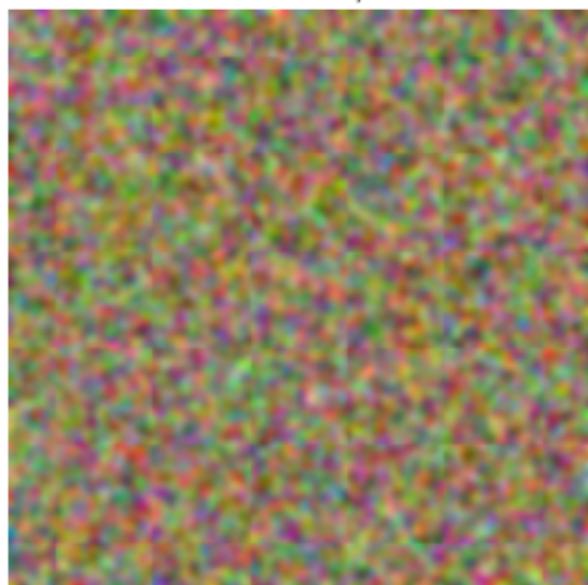
y\_hat:doormat, welcome mat  
y:gorilla, Gorilla gorilla  
Iteration 50 / 1000



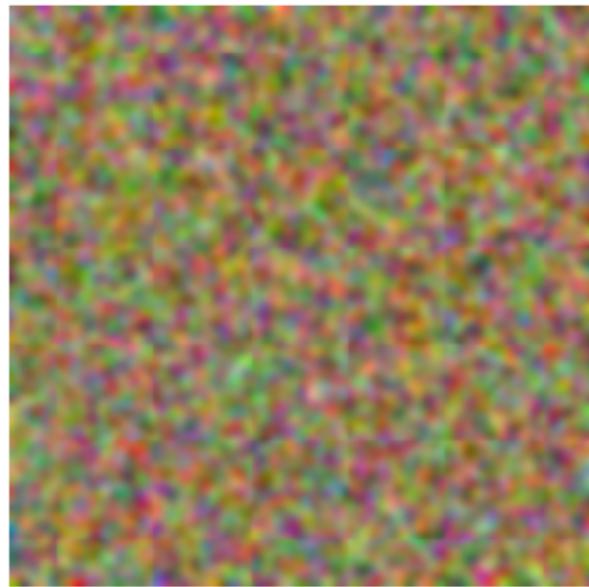
y\_hat: doormat, welcome mat  
y: gorilla, Gorilla gorilla  
Iteration 100 / 1000



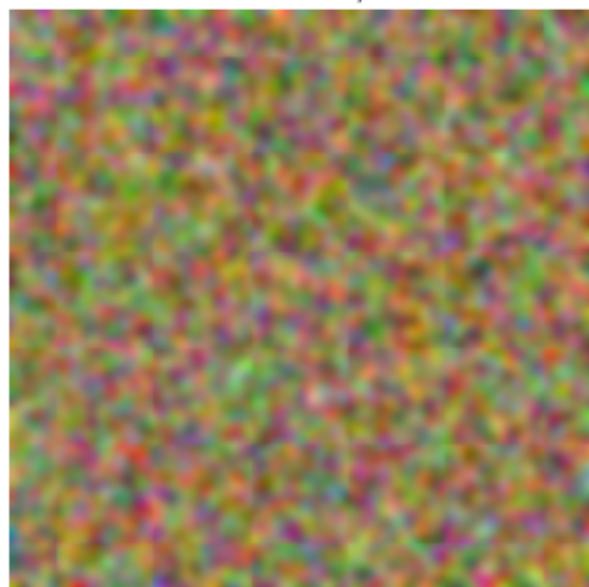
y\_hat: leafhopper  
y: gorilla, Gorilla gorilla  
Iteration 150 / 1000



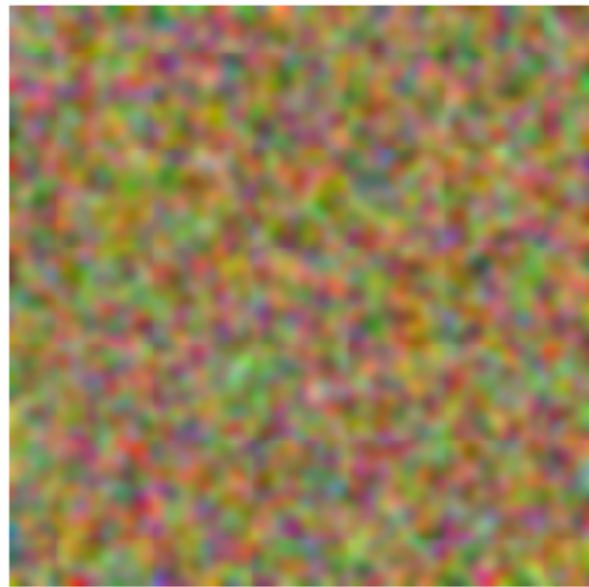
y\_hat:golf ball  
y:gorilla, Gorilla gorilla  
Iteration 200 / 1000



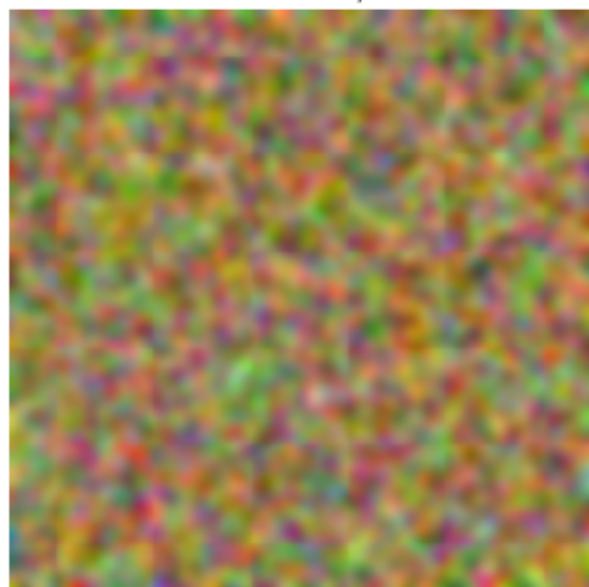
y\_hat:papillon  
y:gorilla, Gorilla gorilla  
Iteration 250 / 1000



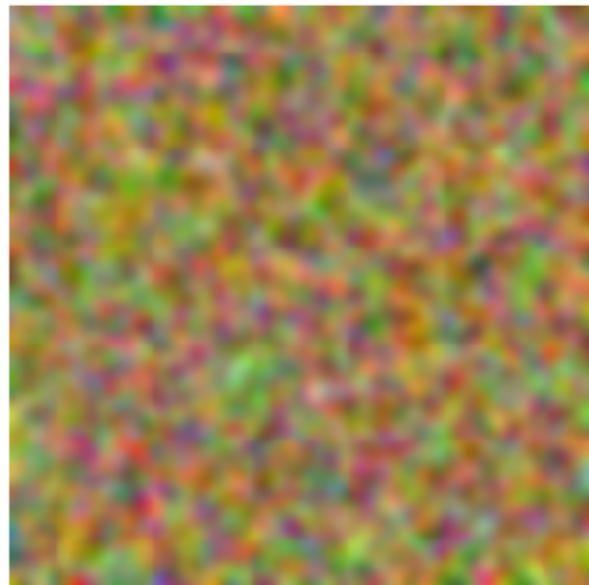
y\_hat:papillon  
y:gorilla, Gorilla gorilla  
Iteration 300 / 1000



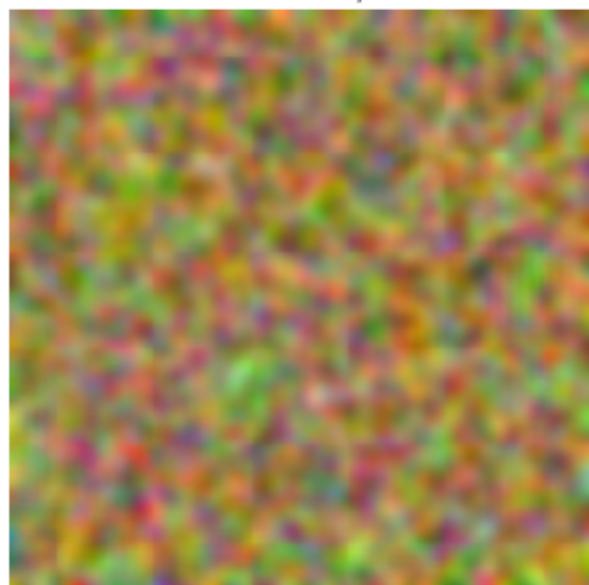
y\_hat:racket, racquet  
y:gorilla, Gorilla gorilla  
Iteration 350 / 1000



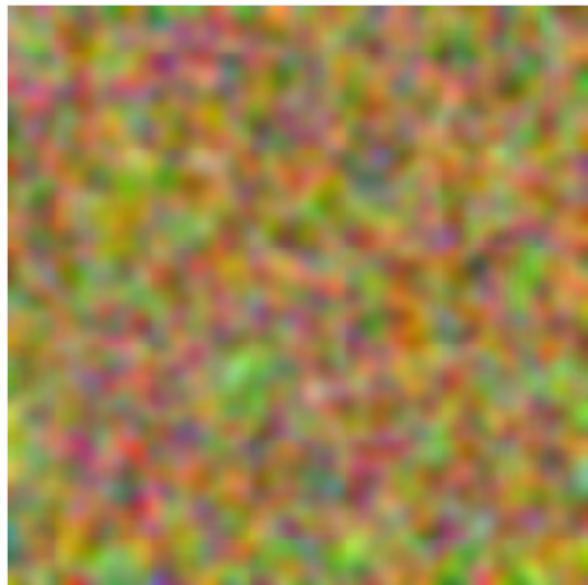
y\_hat:hare  
y:gorilla, Gorilla gorilla  
Iteration 400 / 1000



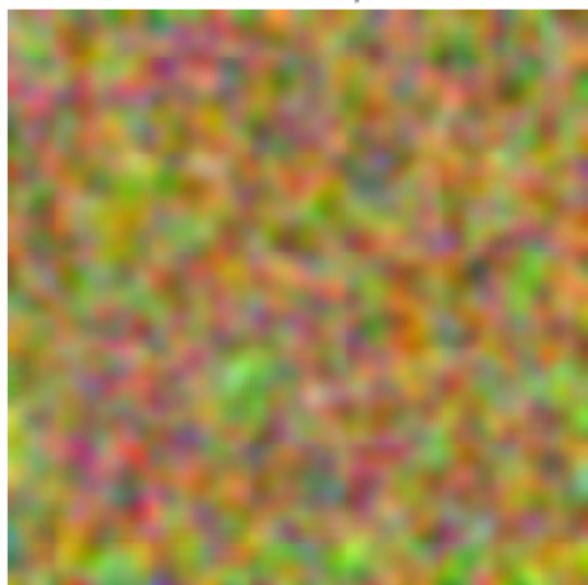
y\_hat:hare  
y:gorilla, Gorilla gorilla  
Iteration 450 / 1000



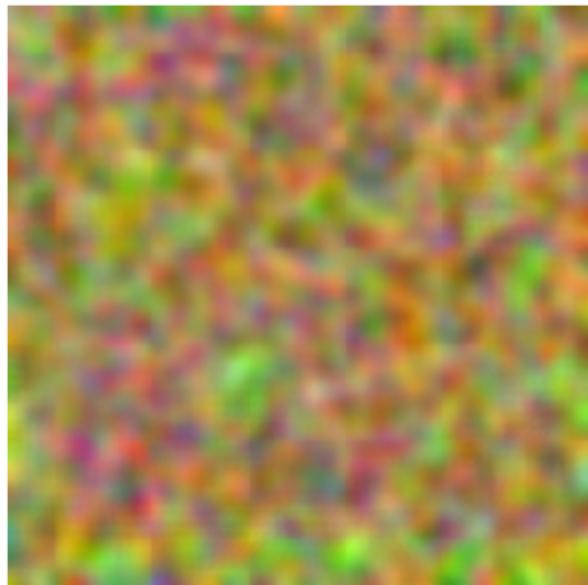
y\_hat: prairie chicken, prairie grouse, prairie fowl  
y: gorilla, Gorilla gorilla  
Iteration 500 / 1000



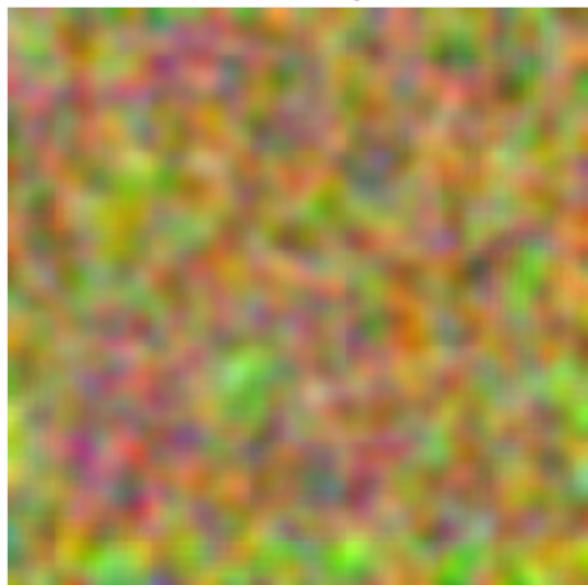
y\_hat: prairie chicken, prairie grouse, prairie fowl  
y: gorilla, Gorilla gorilla  
Iteration 550 / 1000



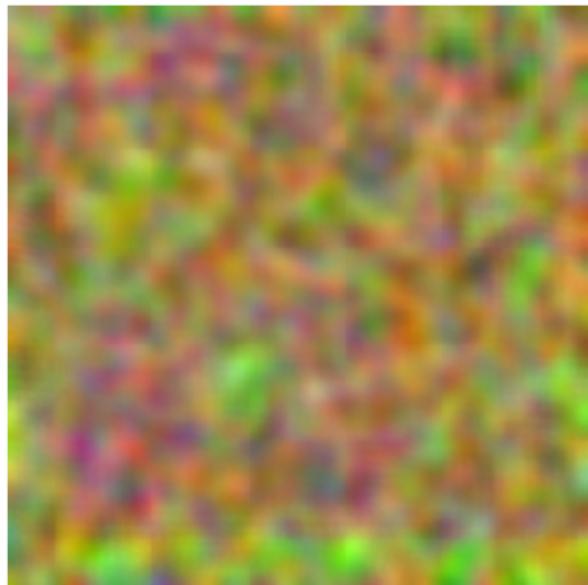
y\_hat: prairie chicken, prairie grouse, prairie fowl  
y: gorilla, Gorilla gorilla  
Iteration 600 / 1000



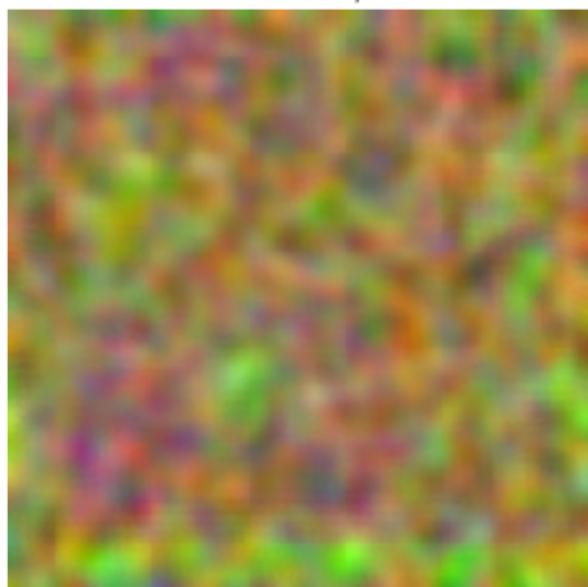
y\_hat: prairie chicken, prairie grouse, prairie fowl  
y: gorilla, Gorilla gorilla  
Iteration 650 / 1000



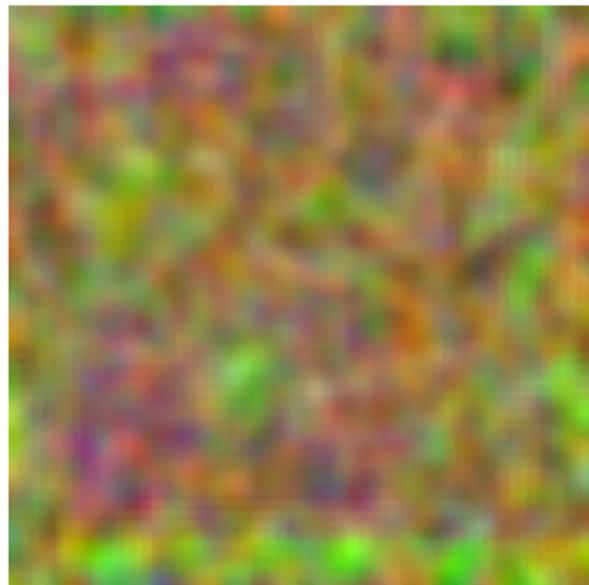
y\_hat: prairie chicken, prairie grouse, prairie fowl  
y: gorilla, Gorilla gorilla  
Iteration 700 / 1000



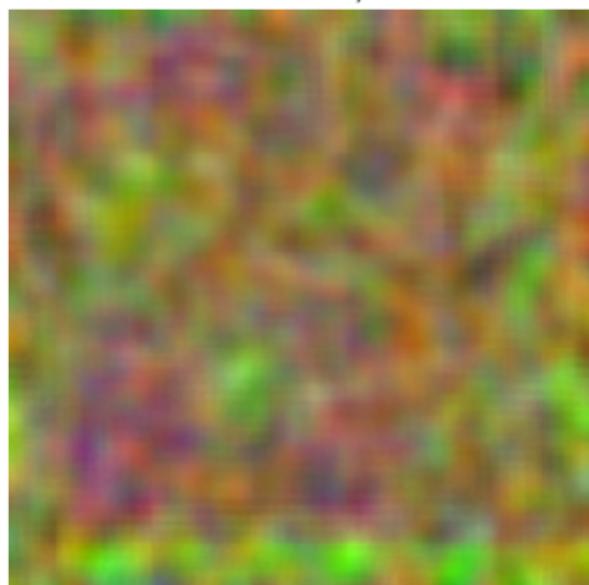
y\_hat: prairie chicken, prairie grouse, prairie fowl  
y: gorilla, Gorilla gorilla  
Iteration 750 / 1000



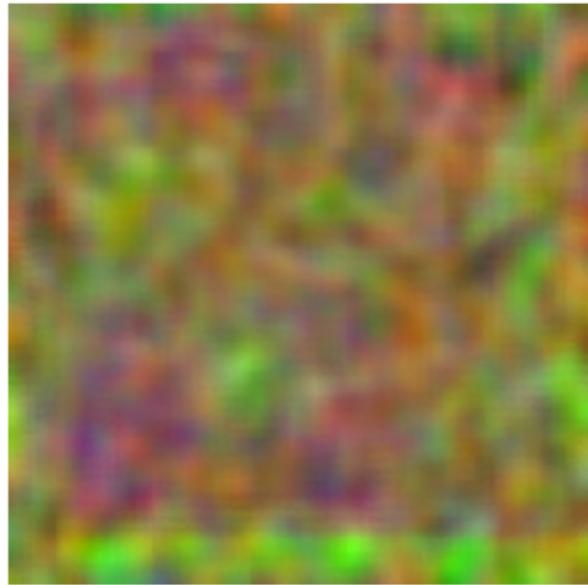
y\_hat: prairie chicken, prairie grouse, prairie fowl  
y: gorilla, Gorilla gorilla  
Iteration 800 / 1000



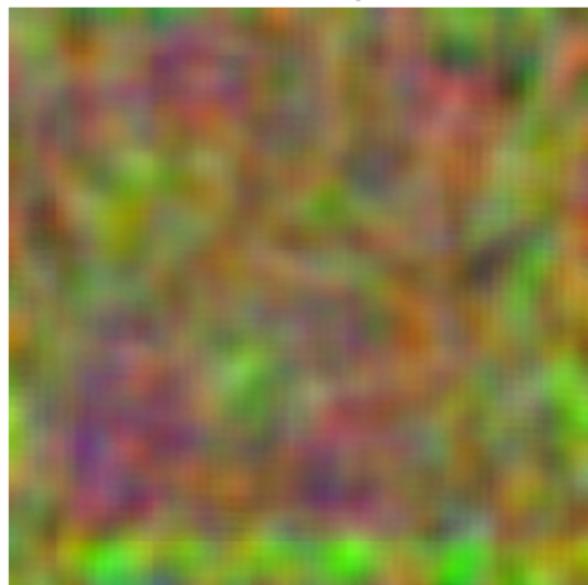
y\_hat: prairie chicken, prairie grouse, prairie fowl  
y: gorilla, Gorilla gorilla  
Iteration 850 / 1000



y\_hat:dhole, Cuon alpinus  
y:gorilla, Gorilla gorilla  
Iteration 900 / 1000



y\_hat:fox squirrel, eastern fox squirrel, Sciurus niger  
y:gorilla, Gorilla gorilla  
Iteration 950 / 1000



y\_hat: fox squirrel, eastern fox squirrel, *Sciurus niger*  
y: gorilla, *Gorilla gorilla*  
Iteration 1000 / 1000

